

Working Without Wages: The Consequences of Widespread Pay Delays*

Dauda Musa
The Spartak Consult

Daniel J. Sonnenstuhl (JMP)
University of Chicago

October 27, 2025

Frequently updated, view latest version [here](#)

Abstract

In this paper, we study a rarely documented firm practice in low-income countries that affects millions of workers and has far-reaching labor market implications: the withholding of employees' wages. Using original survey data from Lagos, Nigeria, we find that 30 percent of workers across firms of all sizes report delayed or unpaid salaries. We develop a theoretical framework showing that, when contractual enforcement is weak, wage withholding can be nearly costless for firms and may even increase employee effort. To examine how wage withholding affects workers' behavior and labor-market participation, we implement a field experiment. We find that delaying wages increases employees' initial effort, without affecting absenteeism or total hours worked. Signaling salary reliability significantly increases job take-up by about 25%, an effect driven primarily by individuals who had initially expressed no interest in wage employment. However, these workers are no more productive than those who accepted jobs without a salary guarantee. Combining intensive- and extensive margin estimates suggests that, in our setting, firms incur minimal productivity losses of 0.2% from engaging in this practice. In contrast, workers place high value on reliable pay, yet weak enforcement and limited worker alternatives create a moral hazard for firms to withhold wages.

*Daniel Sonnenstuhl thanks Greg Lane, Eduardo Montero, and James Robinson for their continuous guidance and support throughout this project and the Ph.D. We also thank Nathan Barker, Chris Blattmann, Christina Brown, Martin Castillo Quintana, Emanuele Colonnelli, Francisco Eslava, Josh Dean, Oeindrila Dube, Steven Durlauf, Yana Gallen, Josh Gottlieb, Amen Jalal, Anne Karing, Erin Kelly, Michael Kremer, Francesco Loiacono, Danielle Nemschoff, Stephen O'Connell, Evans Osabuohien, Raúl Sánchez de la Sierra, Ben Shaver, Shaoda Wang, Chris Woodruff and audiences at the University of Chicago, SIOE, RES PhD Conference, CSAE, DevPEC, AFE Conference, STEG Early Career Workshop and RWI Berlin for helpful comments and discussions. We thank the entire field team for their outstanding work throughout this project. This research is funded by The Weiss Fund for Research in Development Economics, the Development Economics Center and the Becker Friedman Institute, all at the University of Chicago. This document is an output from the research initiative 'Structural Transformation and Economic Growth' (STEG), a programme funded by the Foreign, Commonwealth & Development Office (FCDO), contract reference STEG_LOA_3650_Sonnenstuhl. The views expressed are not necessarily those of FCDO. This research has been approved by the University of Chicago Social and Behavioral Sciences IRB (Protocol Number: IRB24-0989) and the Covenant Health Research Ethics Committee (Protocol Number: CHREC/529/2024) at Covenant University. This experiment was pre-registered at the AEA registry under AEA registry number: AEARCTR-0015132.

1 Introduction

The organization of modern economies rests on a simple exchange: workers supply labor to firms and are paid wages in return. Economic theory generally assumes that wages are reliably paid, as workers can either seek legal redress or retaliate against employers in response to wage disputes (e.g., Bewley, 1998; Krueger and Mas, 2004; Mas, 2006). Yet in many low-income countries, firms frequently delay or entirely withhold wages, undermining this fundamental exchange. In Nigeria, for example, 30 percent of employees have experienced problems receiving their salaries, including substantial delays and partial or complete non-payment.¹ This implies millions of affected workers in Nigeria alone, and similar problems are widely reported in other low-income countries.²

Intuition suggests that worker responses — retaliation against withheld wages and high-ability workers selecting out of wage employment — should impose substantial costs on firms and make delaying salaries unattractive. The prevalence of wage withholding therefore raises a central question: how do workers in low-income countries respond to withheld pay, and how does this practice shape labor-market participation and workforce composition? Studying these margins, however, poses empirical challenges. Measuring workers’ on-the-job responses requires linking performance data to rarely available information on delayed or unpaid wages. In addition, identifying how wage withholding shapes selection into wage employment requires observing how workers deterred from applying to jobs by salary uncertainty respond to job offers once that uncertainty is alleviated.

In this paper, we address these challenges by combining original survey data, a theoretical model, and a randomized controlled trial (RCT) to study the largely unexplored practice of wage withholding. Drawing on the survey evidence we collected in Lagos, Nigeria, we first document the prevalence of wage withholding: 30 percent of workers report delayed

¹Authors’ calculations from original survey data; Section 2.2 discusses these numbers in more detail.

²Extensive anecdotal evidence exist across Sub-Saharan Africa, for example in the form of newspaper articles (Fernandes-Brough, 2023; Akuopha, 2023; Punch, 2023).

or unpaid salaries, with incidents occurring across all firm sizes. To interpret these patterns and generate predictions about worker behavior, we develop a model in which wage withholding can influence effort, job acceptance, and workforce composition. We then test the model’s predictions in a field experiment in Lagos. We find that delaying wages modestly increases employee effort rather than eliciting retaliatory behavior. However, signaling salary reliability raises job take-up, especially among individuals otherwise disinterested in wage employment. Yet, these workers are indistinguishable from typical employees in both productivity and observable characteristics, suggesting that firms face minimal productivity losses from engaging in wage withholding.

To document workers’ experiences with wage arrears, we collected original survey data from 1,279 employees and self-employed individuals.³ We document six new facts about wage withholding, showing that wage arrears are widespread and economically meaningful. First, 30 percent of respondents report withheld or unpaid wages, and second, the median delay is one month. Third, wage withholding occurs across firms of all sizes. Fourth, workers rarely take action in response — only 19 percent left their job, and only one percent pursued recovery through legal or union channels. Fifth, these experiences directly translate into widespread concerns about employers’ reliability: 47 (31) percent of respondents worry that small (large) employers might fail to pay wages as agreed. Finally, employees attribute withheld wages to both business difficulties and deliberate withholding by employers.

After documenting the prevalence of wage arrears, we turn to the question of why this practice arises. We do so by analyzing workers’ responses to wage arrears, which shape employer incentives and payment behavior. We develop a model in which worker reactions to delayed payments play a central role. Specifically, we model wages as non-binding promises that may or may not be paid each period. Unpaid wages accumulate as outstanding balances, rolling over into subsequent periods. Recovery can occur in two ways. First, workers are forward-

³Sampling employees as well as self-employed individuals we specifically account for the fact that some people might have left the labor force because of adverse experiences in previous employment, for example in relation to salary payments.

looking and form beliefs about the likelihood that firms will repay outstanding balances — effectively adding arrears to the wages owed — if the employment relationship survives. The survival, in turn, depends on workers’ effort, following the logic of Lazear (1979) and Huck et al. (2011). Second, if the worker leaves the firm, repayment depends on the institutional environment, such as the quality of the legal system.

We use our model to generate three key predictions. First, the relationship between wage arrears and optimal worker effort — our first empirical target parameter — is ambiguous and depends on the institutional environment. The model predicts that in settings such as Nigeria, where payment uncertainty is high, unpaid wages may initially incentivize increased worker effort. Second, when uncertainty about wage payment is high, increasing payment certainty raises workers’ willingness to accept wage employment — our second empirical target parameter. Third, whether this selection effect matters for the composition of the workforce — in terms of worker productivity, our third empirical target parameter — is ambiguous because selection can be driven by higher ability or more risk averse workers.

To test our model’s predictions and estimate the three target parameters, we conducted a field experiment in Lagos. We incorporated a local firm and recruited individuals for in-person image-labeling tasks tailored to this experiment to precisely measure productivity. The design followed a two-stage randomization. In the first stage, we recruited individuals for our newly created jobs through physical advertisements targeting active jobseekers ($n = 638$) and through in-person recruitment of individuals not seeking wage employment ($n = 1,079$).⁴ Recruits were then assigned to one of three employment conditions varying in salary certainty: (i) high-certainty, guaranteeing timely wage payments; (ii) low-certainty, explicitly stating a 10 to 50 percent chance of non-payment in a given pay cycle; and (iii) a control condition with no payment information. All recruits were required to signal their willingness to work under their assigned terms by attending an in-person orientation day.⁵

⁴All recruitment was conducted in collaboration with a local recruitment agency and recruiters presented themselves as working for a third party recruitment company working on behalf of a client.

⁵Every recruit received a letter outlining their terms of potential employment. While the stated terms of

Those who attended formed the pool from which we hired subsets of candidates for the second stage. We employed 600 candidates across two employment rounds, with a varying composition of the recruited sample. Among those whose terms explicitly mentioned possible payment delays, we randomly assigned salary delays during employment. Treated employees experienced delays lasting one to three biweekly pay cycles over a three-month period.⁶

Our results directly address the three target parameters defined in our conceptual framework. First, we find that being owed salary leads to a modest but statistically significant increase in employee productivity: image classification performance increased by 0.5 percent ($p < 0.01$) on our continuous accuracy index, corresponding to a 5.6 percent ($p < 0.05$) improvement in flawless task completions. These effects are robust across a range of checks and driven by employees with weaker outside options, for whom job loss or forgoing unpaid salaries would be especially costly. Results also remain unchanged for the ex-ante designated “high-stakes” images, where employees were asked to exert greater effort but also had an opportunity to retaliate against the employer. In contrast, absenteeism and hours worked did not change significantly in response to salary delays. This finding aligns closely with our first core prediction, suggesting that in environments with limited legal recourse, outstanding wages can incentivize greater worker effort rather than provoke retaliation.

Next, we test whether salary uncertainty deters individuals with higher outside options from selecting into wage employment. We find that salary certainty matters — particularly for these individuals. Our key finding is that receiving a terms-of-employment letter explicitly conveying higher salary certainty significantly increased job interest, especially among in-person-recruited individuals who initially expressed no interest in the job. Orientation attendance in this group increased by 11 percentage points ($p < 0.01$), corresponding to a 25 percent effect size. In contrast, job acceptance exceeded 95 percent among individuals who actively responded to job advertisements and have measurably lower outside options,

employment were fixed, participants were aware that actual employment was not guaranteed and depended on vacancy availability and their confirmed interest.

⁶Ethical considerations are discussed in Section 4.7.

with no significant variation across treatment arms — consistent with our model’s prediction that salary certainty should matter primarily for those with stronger outside options. We interpret these findings as evidence that pervasive salary uncertainty deters a substantial share of the population from pursuing wage employment.

The implications of this selection for firms depend primarily on worker productivity, as the Nigerian labor market is relatively slack and firms can readily find replacements.⁷ To assess whether selection into employment matters, we examine whether workers attracted by reliable payment differ in productivity or characteristics from typical employees. We first analyze selection into wage employment using a standard LATE framework, treating assignment to the high-certainty employment condition as an instrumental variable for job take-up. Identifying the characteristics of compliers — individuals induced to accept the job by the high-certainty employment condition — shows that they are statistically indistinguishable from always-takers, employees who accept wage employment even in the control condition. Second, we assess productivity differences between the in-person recruited sample and the job-advertisement recruited sample and find that they perform equally well. These results suggest that although reliable payment attracts additional workers, it does not appear to substantially affect the quality or composition of the workforce, indicating only limited productivity implications for firms.

Finally, leveraging randomized salary offers, we quantify individuals’ monetary valuation of salary certainty. We calculate the marginal rate of substitution between higher salary certainty and the monetary salary offer to obtain a willingness-to-pay (WTP) measure.⁸ We find that individuals value higher salary certainty substantially, estimating a WTP of around

⁷The latest official *labor underutilization* rate, reported for 2024, was 14.5% (National Bureau of Statistics, 2024). This measure combines unemployment and time-related underemployment, i.e., individuals who work less than full-time and would like to work more. In a context where many people hold informal jobs that are rarely full-time, this is the most appropriate indicator for assessing labor market slack. As an empirical reference for labor market slackness, we were able to recruit more than 300 employees within roughly three weeks of posting job advertisements for our experiment.

⁸Strictly speaking, this measure captures the amount of salary individuals were willing to forgo for greater salary certainty, but interpreting it as a WTP remains conceptually valid.

25,000 Nigerian Naira (NGN) equivalent to around USD 15 and corresponding to over 100 percent of the median weekly wage. Additionally, we validate these estimates using a choice experiment administered during the initial job interview or information session, which yields strikingly consistent results.

Our findings suggest that environments with weak enforcement and limited worker recourse create firm-side moral hazard. Firms can engage in wage withholding while facing limited risk of legal challenge and incurring only minor productivity losses. A back-of-the-envelope calculation combining our extensive- and intensive-margin estimates indicates that, in our study setting with slack labor markets, firms' productivity decreases by about 0.2 percent. This decline would likely be larger in environments where labor markets are tighter and replacing workers is more costly. At the same time, wage withholding imposes substantial welfare costs on workers, reflected in their high valuation of reliable pay. Taken together, these results suggest that firms face little incentive to refrain from wage withholding, highlighting the potential role of stronger institutional enforcement in improving conditions for workers and discouraging such practices.

One caveat to our findings is that we do not capture potential general-equilibrium effects on firms. In our data, workers who avoid wage employment are more likely to enter self-employment, where they compete with firms. This competition may impose additional costs on firms — an important avenue for future research. A further limitation is that, although wage withholding directly reduces workers' welfare, it may also generate offsetting effects that we cannot measure. Firms might, for instance, be less willing to hire if withholding wages were not an option.

Prior to this study, wage arrears in low-income countries have been documented in public-sector contexts (Diamond and Schiller, 1993; Flynn and Pessoa, 2014; Buehren et al., 2018). More broadly, documentation and analysis through a macroeconomic lens exist for post-Soviet Russia (Alfandari and Schaffer, 1996; Clarke, 1998; Lehmann et al., 1999; Earle and

Sabirianova, 2002; Gerry et al., 2004; Lehmann and Wadsworth, 2007; Earle and Peter, 2009) and Ukraine (Boyarchuk et al., 2005). While widespread wage arrears are typically not an issue in industrialized economies, related phenomena such as wage theft — especially minimum-wage violations — are nevertheless well-documented in low-skilled labor markets in these contexts (Bernhardt et al., 2009; Milkman et al., 2010; Robinson et al., 2011; Galvin, 2016; Clemens and Strain, 2022). However, to the best of our knowledge, no prior research has explicitly investigated how outstanding wage balances — either individually or in aggregate — affect worker effort and selection into employment.

Our analysis builds on existing research that examines worker performance and behavior, especially in low-income countries (e.g. Falk, 2007; Kaur et al., 2015; Breza et al., 2018; Freeman et al., 2025; Kaur et al., 2025). We demonstrate, theoretically and empirically, that salary delays can serve as an incentive mechanism by motivating workers to increase effort to secure future payments. This finding aligns closely with evidence showing that performance-linked incentives increase worker productivity in low-income contexts (Bandiera et al., 2007; Ashraf et al., 2014, 2018; Guiteras and Jack, 2018) and more generally (e.g. Lazear, 2000).

We also contribute to the literature on labor-market frictions in low-income countries (see Breza and Kaur (2025) for a recent overview). A growing body of evidence documents worker reluctance toward wage employment, driven by preferences for flexibility (Blattman and Dercon, 2018) — potentially linked to habit formation (Cefala et al., 2024) — complementary labor supply (Donald and Grosset, 2024), redistributive pressures (Carranza et al., 2022), and cultural norms (Cassan et al., 2022; Oh, 2023). Kapoor (2025) investigates how liquidity constraints and limited firm commitment shape contractual preferences among job-seeking day laborers in Indian spot-labor markets. We go beyond this evidence to show that salary uncertainty — arising broadly from widespread wage arrears — constitutes a general and significant deterrent to wage employment. More broadly, our findings highlight that payment uncertainty shapes individual behavior, consistent with evidence from other contexts (Dunn,

Gottlieb, Shapiro, Sonnenstuhl and Tebaldi, 2024).

Finally, we contribute to existing research documenting the economic impacts of weak institutional enforcement in low-income countries. Prior work shows that limited enforcement contributes to inefficient firm behavior, including inferior management practices (Laeven and Woodruff, 2007; Bloom and Van Reenen, 2010; Bloom et al., 2013; McKenzie and Woodruff, 2017), and resource misallocation (Boehm and Oberfeld, 2020), thereby constraining firm size and reducing overall productivity (Akcigit et al., 2021). We extend this literature by explicitly documenting how firms can strategically leverage weak enforcement via wage withholding to incentivize worker productivity. In doing so, we introduce a novel margin to a broad literature that examines a plethora of dimensions through which limited contract enforcement influences economic outcomes (e.g. Djankov et al., 2003; Chemin, 2012; Sánchez de la Sierra, 2021; Sánchez de la Sierra et al., 2024).

2 Salary Uncertainty in Nigeria: Context and Empirical Patterns

In this section, we briefly outline the broader context of our study and present novel descriptive evidence on salary uncertainty and wage withholding in urban Nigeria. We document six new facts that describe its scale, incidence, and workers’ responses. The institutional and economic context is essential to conceptualize both the emergence and consequences of widespread wage arrears. Although our data and contextual description are specific to Nigeria, the setting shares key features with labor markets across many low-income countries.

2.1 Context

The setting for this study is Lagos, a megacity in southwestern Nigeria with approximately 18 million inhabitants. Lagos is Nigeria’s largest city and economic hub, although Nigeria remains among the poorest countries in the world, with a GDP per capita of \$1,596.9 in 2023 (World Bank, 2023 USD). Nigeria’s economy is heavily dependent on oil, making it sensitive to fluctuations in global oil prices and causing recurrent inflationary pressures.

This volatility exacerbates financial insecurity for firms and workers alike. While the global technology boom has also reached Nigeria, making Lagos home to one of Africa’s most vibrant startup and tech sectors, the majority of people still work in more traditional economic activities. Retail and manufacturing, for instance, are important sectors both in Lagos and across Nigeria (PwC, 2024). Many of these businesses are notably small and clustered in market-like settings; as an example, Lagos’s *Computer Village* is a market named after its concentration of small-scale IT and phone retailers. Most of these shops are independently owned, typically employing only a few individuals. In fact, micro, small, and medium-sized enterprises collectively employ around 80 percent of Nigeria’s workforce and generate over 40 percent of its GDP (PwC, 2020, 2024).⁹ Micro enterprises constitute over 95 percent of these businesses, making the median enterprise in Nigeria effectively a one-person operation. The vast majority of these enterprises operate informally, with limited access to formal financial systems and minimal regulatory oversight.

Lack of effective governance and regulatory enforcement also characterizes Nigeria’s broader institutional environment. A recent U.S. Department of State Human Rights Report describes widespread corruption, substantial shortcomings in enforcing employment laws — violations of minimum wage regulations and basic labor standards are rarely investigated — and an overall grim human rights situation in Nigeria (US Department of State, Bureau of Democracy, Human Rights and Labor, 2023). This characterization resembles scholarly assessments of Nigeria’s rule-of-law environment, which document compromised judicial independence, political interference, and routine disregard for judicial decisions (e.g. Akomolafe, 2021; Igwe, 2021). Nigeria’s consistently low global position in the World Justice Project’s rule-of-law index reflects these conditions.¹⁰ Such conditions are not unique to Nigeria; un-

⁹Micro enterprises generally have 10 or fewer employees, small enterprises 11–50 employees, and medium-sized enterprises 51–200 employees.

¹⁰The World Justice Project is an independent organization that publishes annual global rule-of-law rankings based on factors such as constraints on government power, absence of corruption, regulatory enforcement, and protection of fundamental rights. Nigeria has consistently ranked 120th or lower out of 142 countries over the past three years.

reliable law enforcement and the lack of an impartial, well-functioning judiciary are common institutional characteristics in low-income countries (Sánchez de la Sierra, 2021; Sánchez de la Sierra et al., 2024).

2.2 Survey Evidence: Six New Facts on Wage Withholding

We conducted an original survey to document new empirical facts about wage withholding in Nigeria. Between June and August 2025, we surveyed 1,279 individuals in Lagos. To obtain a representative sample of the city’s low- and medium-skilled labor force, enumerators approached respondents at busy public locations using a randomized skip pattern. Because our focus is on workers’ experiences with wage withholding, we restricted the sample to individuals who reported current or previous employment. The main survey findings are presented in Figure 1, which summarizes six new facts about wage withholding.


Fact 1: Wage withholding is prevalent. Panel (a) shows that salary delays and non-payments are widespread. Overall, 29.9 percent of survey respondents reported having experienced at least one form of salary difficulty. We define salary difficulty as experiencing one or more of the following: (i) delayed salary payments, (ii) partial salary payments, or (iii) complete non-payment of wages. Among respondents, seven percent reported receiving only partial payments and four percent reported no payment at all.¹¹ The most common problem, however, was delayed salary payments, affecting 19 percent of respondents.¹²

Fact 2: The median delay is one month. Panel (b) shows the distribution of delay durations. Among respondents who experienced payment delays, the median duration of their longest reported delay was one month. About 25 percent of affected workers experienced delays of four months or longer.

¹¹ Respondents could select more than one category, but only very few did.

¹² Appendix Table A.5 shows that individuals expect wage withholding to occur with a very high probability if they were to start working at a firm tomorrow.

Figure 1: Facts About Wage Withholding



NOTE: This figure presents the six key facts established by our descriptive survey. Panel (a) shows the overall prevalence of wage withholding, and Panel (b) displays the distribution of delay durations. Panel (c) reports the occurrence of wage withholding by firm size, while Panel (d) illustrates employees' reported responses to wage withholding. Panel (e) summarizes workers' concerns about wage employment, and Panel (f) documents their perceptions of why wages were withheld.

Fact 3: Wage withholding occurs across firms of all sizes. Panel (c) shows that delayed and unpaid wages occur across firms of all sizes. We group firms into three size categories. Salary difficulties become less common as firm size increases: 40 percent of respondents reported their worst salary difficulty at small firms (fewer than 10 employees), 32 percent at medium-sized firms (10 – 50 employees), and 19 percent at large firms (more than 50 employees). While a large share of cases remains concentrated in small firms, there is still a substantial probability of experiencing salary difficulties even in very large and formal firms. This pattern underscores that wage withholding is not limited to the informal sector or small firms but affects the entire economy.¹³

Fact 4: Worker responses are limited. Panel (d) shows how employees respond to wage withholding. Only 19 percent of employees reported leaving their jobs in response to not receiving their salary — a surprisingly low share given that paying the agreed wage is among an employer’s core obligations to employees.¹⁴ Moreover, firms appear to face minimal legal risk, as almost no respondents reported taking legal action or seeking union assistance.¹⁵ Nevertheless, employees largely disapproved of salary withholding; only eight percent indicated understanding of the situation.¹⁶

Fact 5: Wage withholding is a substantial concern. The widespread occurrence of wage withholding translates into strong concerns about being paid as agreed when considering employment. Panel (e) summarizes these concerns, showing that reliability of pay is the primary issue workers associate with both large and small firms.¹⁷ Just over 47 percent

¹³Panel B of Appendix Table A.4 shows that 49 percent of survey respondents know that someone within their direct social environment has experienced salary difficulties.

¹⁴Panel A of Appendix Table A.3 provides details on how respondents coped financially with unpaid salaries.

¹⁵Public sector employees also experience wage withholding for example.

¹⁶In Panel A of Appendix Table A.4, we provide additional evidence that the social norm and clear expectation is for firms to pay their employees fully and on schedule.

¹⁷Survey participants were asked whether they had any concerns when thinking about employment and, if so, which ones. To avoid biasing responses, they could express more than one concern, so percentages do not sum to 100 but reflect the frequency with which each concern was mentioned.

of respondents expressed concern about salary payments when thinking about employment at small informal firms, and 31 percent expressed concern about employment at large formal firms.¹⁸ In addition, respondents also expressed concerns about employers not honoring other aspects of the work agreement (e.g., working hours), job loss, and religious or ethnic discrimination in the workplace.¹⁹

Fact 6: Workers perceive wage withholding as both involuntary and deliberate.

Panel (f) illustrates employees’ perceptions of why their wages were withheld. Employees attribute a substantial share of these difficulties either to employers’ inability to pay due to poor business conditions (40 percent) or to deliberate wage withholding (27 percent). These responses suggest two main perceived reasons for wage withholding: employers’ liquidity constraints and intentional nonpayment.

3 Theoretical Framework

In this section, we develop a theoretical framework to understand the mechanisms driving workers’ responses to unpaid wages. The purpose of the model is twofold. First, it illustrates the trade-offs workers face when their wages remain unpaid. Second, it derives predictions about worker behavior that guide our empirical analysis.

We model wages as non-binding promises and study the decision problem of a forward-looking worker who must choose how much costly effort to exert when wages remain unpaid. Unpaid balances roll over into the next period and are added to the amount the worker *ought to* receive. When accepting the job, the worker holds an initial belief that the firm will pay wages as agreed. As unpaid balances accumulate, the worker updates these beliefs about the firm’s future payments and compares the continuation value from remaining at the firm with the

¹⁸Because employment in small businesses owned by members of one’s social network is common, we also present results from a randomized subset of respondents who were explicitly asked to consider employment at a small business owned by a family member or friend. The level of concern remains largely unchanged.

¹⁹In Appendix Table A.6 we provide additional descriptive evidence that salary uncertainty influences stated employment preferences. Our survey respondents overwhelmingly express a preference for self-employment over wage employment mirroring synthesized findings by Breza and Kaur (2025).

value of their outside option. The effort decision therefore depends on the worker’s evolving beliefs and the attractiveness of alternatives outside the firm. The framework focuses on worker behavior in an environment characterized by weak enforcement and uncertain wage payments. To capture this environment in the simplest form, we assume firms differ only in their payment reliability: a strategic firm that never pays and a non-strategic firm that pays in full unless hit by a liquidity shock. This reduced-form setting generates a simple belief structure that links workers’ expectations and outside options to their effort responses. We match the model to empirically elicited beliefs about payment probabilities, which allows us to derive qualitative predictions for the context we study.

In Appendix F, we extend the framework by endogenizing firms’ repayment behavior and show that the simplified version captures the essential worker-side mechanisms without loss of generality. We then numerically solve for an equilibrium in which workers’ beliefs and firms’ repayment policies are jointly consistent.

Setup. Time is discrete and firms are of two permanent types: strategic (S) and non-strategic (N). Firms offer wage w to workers, but strategic firms never pay, whereas non-strategic firms pay unless experiencing liquidity shocks. Each period, non-strategic firms face i.i.d. liquidity shocks $L \sim \text{Geom}(\rho)$ that end with probability ρ implying that a shock can persist across multiple periods. During a shock, the firm withholds the wage; when the shock resolves, it resumes full payment.

Workers hold priors over firm types: let λ_0 denote the probability that a newly matched worker believes their employer is of type S . In each pay cycle t , firms either pays the full amount due or nothing — if they are the strategic type or face a liquidity shock. Let $\chi_t \in \{0, 1\}$ denote the realized payment decision applied to wages outstanding from past periods and the current period’s wage, so that arrears evolve according to $B_{t+1} = (1 - \chi_t)(B_t + w)$.

Workers observe χ_t , update their beliefs by Bayes’ rule about the firm type they face, and

thus the likelihood of future payment, which evolves according to

$$\lambda_t(\lambda_0, \rho) = 1 - \frac{(1 - \lambda_0) \exp(-\rho B_t/w)}{\lambda_0 + (1 - \lambda_0) \exp(-\rho B_t/w)}, \quad \lambda_0 \in (0, 1), \rho > 0 \quad (1)$$

for as long as workers remain unpaid. After updating beliefs, workers choose costly effort e_t which determines the probability $p(e_t)$ that the match continues ($p'(e_t) > 0$ and $p''(e_t) < 0$). If the match breaks, the firm must repay outstanding wages, scaled by a penalty factor $\xi > 1$, with probability ϕ .

3.1 The Worker's Problem

The worker observes B_{t+1} — the amount of owed wages carried into the next pay cycle — after the realized payment χ_t and then chooses effort e_t . Effort affects both current utility, through its convex cost $\psi(e_t)$ and the continuation value, since higher effort raises the probability that the match continues into the next period, reflecting Lazear (2000). Future utility depends on whether the match survives or breaks. If the match continues, utility reflects the expectation of future payments after workers update their beliefs about the firm type they face; if the match ends, it reflects the outside option and the probability of recovering outstanding wages. The worker's value function is therefore

$$\begin{aligned} V(B_t; \lambda_0, \rho) = \max_{e_t} \Big\{ & \overbrace{u(\chi_t(B_t + w))}^{\text{observed payment}} - \psi(e_t) \\ & + \beta \left[p(e_t) \mathbb{E}[V(B_{t+1}; \lambda_0, \rho)] + (1 - p(e_t)) (V^{\text{out}} + \phi u(\xi B_{t+1})) \right] \Big\}. \end{aligned} \quad (2)$$

Equation (2) shows that the worker's value depends on two key objects corresponding to the two possible outcomes of the match. When the match breaks with probability $1 - p(e_t)$, the worker receives the outside option and, with probability ϕ , any recovered arrears and penalties. We treat ϕ as exogenous, capturing the strength of the institutional environment. In settings with well-functioning judicial systems, ϕ may approach one, reflecting a high

probability that workers can successfully recover arrears. In contrast, in environments such as Nigeria, where enforcement mechanisms are weak, ϕ is likely close to zero, reflecting the limited ability of workers to reclaim unpaid wages.

When the match continues with probability $p(e_t)$, the worker receives the expected continuation value, $\mathbb{E}[V(B_{t+1}; \lambda_0, \rho)]$, which captures the tension workers face. As arrears accumulate, the total amount owed increases, mechanically raising the expected utility from remaining with the firm since more pay is due in the next period. However, beliefs about the likelihood of receiving payments in the future may counteract this effect. The belief of facing a firm that eventually repays is downward sloping in B and hence higher arrears have an ambiguous effect on the worker’s expected continuation value.

3.2 Predictions of the Model.

We now use the model to derive three predictions about worker behavior and productivity. To derive qualitative predictions for the setting we are studying, we estimate λ_0 and ρ via nonlinear least squares using workers’ beliefs about the likelihood of eventual payment after sustained nonpayment, elicited in the survey described in Section 2.2.²⁰

Qualitatively, the context of our study is characterized by three key features. First, there is a low probability of recovering unpaid wages through the legal system (see Figure 1d). Second, workers hold a high prior belief that they may face a firm that does not pay wages as agreed (see Appendix Table A.5). Third, workers maintain relatively strong beliefs that they will eventually be paid even after experiencing prolonged nonpayment (see Appendix Figure xxx). These features are captured by the parameterization $\lambda_0 = 0.41$, $\rho = 0.14$, and $\phi = 0.2$.²¹

²⁰See Appendix F for details on estimation.

²¹See Appendix F for additional details on model parameterization.

Prediction 1 (Ambiguous Initial Effort Response). Optimal worker behavior hinges on the worker’s first-order condition, which reads

$$\psi'(e_t^*(B_{t+1})) = \beta p'(e_t^*(B_{t+1})) [\mathbb{E}[V(B_{t+1}; \lambda_0, \rho)] - V^{\text{out}} - \phi u(\xi B_{t+1})]. \quad (3)$$

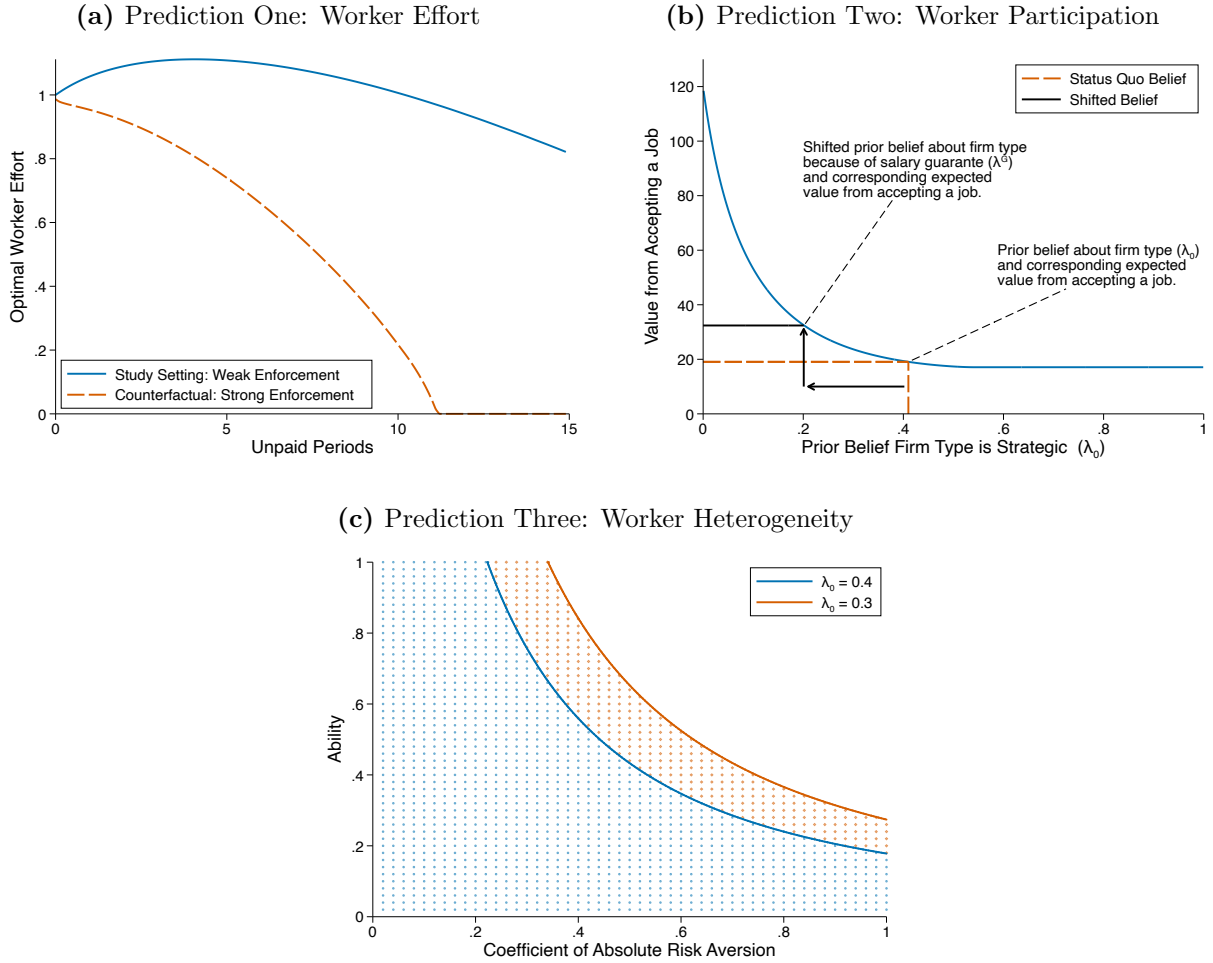
Differentiating the first-order condition with respect to wage arrears B , provides an expression for how optimal effort responds to increases in arrears. Evaluating this expression for $B = 0$, we obtain

$$\left. \frac{de^*}{dB} \right|_{B=0} \geq 0 \iff \mathbb{E}[V_w(w; \lambda_0, \rho)] \geq \phi \xi u'(w). \quad (4)$$

Prediction 1 highlights that the initial effect of wage arrears on employee effort — our first empirical target parameter — is theoretically ambiguous, depending on the enforcement probability ϕ and the marginal value of an outstanding wage balance. This ambiguity is illustrated in Panel (a) of Figure 2. Under a parameterization resembling our study setting ($\phi = 0.02$) effort initially increases and gradually declines as wages remains unpaid. In contrast, when the probability of recovering unpaid wages is high even after the match ends ($\phi = 0.9$), effort declines immediately and much more sharply.

Prediction 2 (Worker Participation). Workers’ beliefs do not only affect their behavior on the job, but also matter for the decision whether to accept a job. Assume a worker accepts employment if and only if $\mathbb{E}[V(B_0; \lambda_0, \rho)] \geq V^{\text{out}}$, i.e. if the ex-ante value of the job exceeds the outside option. Suppose firms can issue a credible signal G , such as a third-party-verified salary guarantee, which shifts workers’ initial prior about firm type to $\lambda_0^G < \lambda_0$ while keeping the remaining environment constant. This guarantee would increase the expected value of the jobs $\mathbb{E}[V(B_0; \lambda_0^G, \rho)] > \mathbb{E}[V(B_0; \lambda_0, \rho)]$ as Panel (b) of Figure 2 illustrates. This implies that take up of jobs should unambiguously increase if the job offer contains a credible signal about the firm type. This is our second empirical target parameter.

Figure 2: Predictions of the Model



NOTE: This figure plots the three predictions of the model. Panel (a) shows workers' optimal effort when wage payments fail to materialize. The blue solid line illustrates the predicted dynamics under weak contract enforcement, while the dashed orange line represents the case with stronger enforcement. Panel (b) shows how the value from accepting a job varies with the prior belief that the firm will never pay the promised wage (i.e. is of type S). Ex ante, this value declines with higher prior beliefs of facing a non-paying firm; thus, any reduction in the initial prior λ_0 increases the expected value of the job. Panel (c) shows the set of workers — characterized by ability and absolute risk aversion — who would accept the risky job over their outside option. The blue solid line depicts workers accepting the job under the prior $\lambda_0 = 0.41$. Any reduction in the belief of facing a nonpaying firm makes wage employment more attractive. Importantly, marginal entrants can differ along two dimensions: higher ability or greater risk aversion. The orange line illustrates this margin by showing the combinations of worker types that would accept wage employment following a shift to $\lambda_0 = 0.3$.

Prediction 3 (Worker Heterogeneity). We now introduce worker heterogeneity along two dimensions: ability and risk aversion. We assume that workers’ outside options increase with ability. Accepting a job is risky when ex-ante uncertainty about wage payments is high, compared to the risk-free but lower income stream available in the outside option. We simulate this trade-off in Panel (b) of Figure 2 showing the combinations of risk aversion and ability for which accepting wage employment is profitable. This exercise highlights a key insight: when the initial belief about firm type (λ_0) decreases, the additional workers who now find wage employment attractive may be either more able or more risk-tolerant. Characterizing the marginal worker is our third target parameter.

4 Experimental Design and Implementation

Our experiment is designed to estimate the three empirical target parameters outlined in the preceding Section 3 by assessing three complementary aspects of salary uncertainty: (i) employees’ on-the-job effort responses to delayed salary payments, (ii) the overall impact of salary uncertainty on labor force participation, and (iii) its effect on workforce composition, specifically regarding the types of workers who choose to accept employment.

To examine these outcomes, we incorporated a firm in Nigeria, The Spartak Consult, whose primary business activity is data classification and labeling. As a newly established firm, The Spartak Consult had no preexisting reputation that could influence employee perceptions or behavior. Through this firm, we were able to extend job offers and hire jobseekers for short-term positions while maintaining complete control over working conditions and salary payments.

Our experimental design follows a two-stage randomization. First, we designed three distinct job offers explicitly varying in terms of salary certainty — these offers form our initial treatment arms. To credibly convey differences in salary certainty to jobseekers, we collaborated with a local recruitment agency so that all salary-related information came from a

third party rather than directly from the employing firm. Through this collaboration, we recruited individuals through two distinct strategies to reach different populations: individuals proactively responding to job advertisements, and self-employed individuals who may not actively seek wage employment. This strategy allows us to assess heterogeneous responses to salary uncertainty across these populations.

In the second stage, conducted in two rounds with varying compositions of the two recruited populations, we hired subsets of interested respondents. We then randomly implement salary delays among respondents who accepted job offers explicitly mentioning the possibility of delayed payments. To precisely measure productivity responses to these treatments, we developed a job task tailored to this experiment.

The experiment ran from January to October 2025. We posted job advertisements from January to early March and received most responses then, though some continued through early June. In-person recruitment took place in two waves, from January to March and again from June to July. Employment occurred in two rounds: February to May, and August to October. End-of-employment surveys indicate that employees perceived the setup as genuine. In both employment rounds, approximately 80 percent stated that image labeling was the primary purpose of their employment arrangement (see Appendix G.1 for details).

4.1 Job Offer Treatments

Employment Terms Treatment 1 (Control Arm). In the first treatment arm, jobseekers are informed that they work for a local Nigerian firm and receive a fixed monthly salary. No additional information is given to jobseekers about payment modalities of salaries or the firm.

Employment Terms Treatment 2 (Uncertainty Arm). In the second treatment arm, jobseekers received the same information as in treatment arm 1, but were additionally told that each pay cycle carried a probability of salary nonpayment due to Nigeria’s difficult

economic conditions. This probability was randomized to take values of 10%, 20%, 30%, 40%, or 50%.

Discussing this risk ex ante with jobseekers is a limitation of the design, as wage withholding typically occurs unexpectedly. Disclosure, however, was necessary for ethical reasons. To preserve uncertainty, jobseekers were informed only that payment might be withheld; hence, if wages were not paid on the scheduled date, employees did not know whether repayment would follow.

Given the adverse economic context, jobseekers were generally unsurprised by the possibility of nonpayment. As shown in Section G.2, receiving this information did not significantly alter individuals' beliefs about timely salary payments.²²

Employment Terms Treatment 3 (Salary Certainty Arm). In the third treatment arm, jobseekers receive the same information as in Treatment Arm 1. Jobseekers are additionally informed that their salaries are guaranteed to be paid on time. The firm uses a third-party automated payment system, directly connected to a bank account that holds sufficient funds to fully cover salaries for the entire duration of the employment period. There would be no uncertainty regarding timely salary payments. Jobseekers are further reassured that no previous employees receiving this offer have reported delayed or unpaid salaries. This treatment provides the credible salary-certainty shock required to identify our target parameter related to workforce participation and composition (target parameter 2 and 3).

4.2 Recruitment and Measure of Take Up

To effectively conduct the recruitment according to our experimental needs, we collaborated with a local recruitment agency, *Unlocking Creativity*. The collaboration allowed some of our enumerators (referred to as recruiters in the following) to be temporarily affiliated with *Unlocking Creativity*, visibly representing the agency during recruitment (e.g., wearing offi-

²²We interpret this as further evidence that wage withholding is sufficiently common that making the possibility explicit does not meaningfully change expectations about being paid on time.

cial identification badges and using agency-branded materials). Recruiters then presented themselves as recruiting for their client, *Spartak Consult*, our company set up for this experiment. This approach enabled us to credibly conduct in-person recruitment leveraging the credibility of the agency’s credentials while retaining complete control over the recruitment team and their training to ensure maximum compliance to the study protocol. Moreover, this collaboration enhanced the credibility of the guaranteed salary assurances in employment terms treatment arm 3. Recruiters’ visible affiliation with an independent agency made their statements about salary payments more credible than if they had come from the client company’s own recruiters, enabling us to credibly vary salary certainty.²³

Job Advertisement Recruiting. To recruit employees actively seeking employment — essential for estimating effort responses to salary delays among typical employees (target parameter 1), and serving as a comparison group for workforce composition analysis (target parameter 3) — members of the field team posted physical job advertisements across selected areas in Lagos from January to early March 2025.²⁴ The form, content, and placement of these advertisements, shown in Appendix Figure B.1, were chosen to closely resemble local norms.²⁵ The advertisements were printed in black-and-white on letter (A4) size paper and pasted on walls, poles, and similar public surfaces in selected areas within a roughly one-hour commuting radius of the work locations. They provided only limited information, simply mentioning the general nature of the work (data classification tasks) and a salary range (50,000 NGN to 85,000 NGN), and instructed jobseekers to contact *The Spartak Consult* via phone call or WhatsApp message to express their interest. Providing only basic information in the job advertisements also served our experimental design. By limiting initial details, we obtained an initial sample of jobseekers who responded to the job advertisement indepen-

²³Follow-up phone surveys, described in greater detail in Appendix G.2 and Appendix Figure G.10, confirm the effectiveness of this treatment in reducing individuals’ concerns about receiving their agreed salary payments.

²⁴We started posting job advertisements on January 11, 2025. Job advertisements were then posted for the remainder of January.

²⁵Appendix Figure B.2 shows our job advertisement among similar job postings in Lagos.

dently of any subsequent information about salary payment conditions. All jobseekers who responded to the job advertisement were invited to a job interview.

Interviews were conducted starting on January 27, 2025, either at a rented event hall or at the company’s office location.²⁶ During the interviews, we collected jobseekers’ baseline characteristics and administered a shortened version of Raven’s Progressive Matrices, as well as a choice experiment designed to elicit jobseekers’ preferences over different types of job offers.²⁷

At the end of each interview, we provided jobseekers with a letter specifying their terms of potential employment. The terms of employment letters differed in terms of the monthly salary offer — randomized between 55,000 NGN and 85,000 NGN (approximately USD 22 to 34) — and also conveyed different information about salary reliability according to the three treatment arms described in Section 4.1.

In-Person Recruiting. To effectively reach our target sample of individuals who are not necessarily looking for employment — and who thus plausibly have higher outside options, crucial for estimating labor force participation (target parameter 2) and workforce composition effects (target parameter 3) — we implemented an in-person recruitment strategy. Recruiters frequented pre-specified market areas and busy public spaces, approaching individuals according to a predetermined skip and selection pattern.²⁸ Approached individuals were invited to participate in an immediate, approximately 20-minute job information session. For individuals who immediately expressed interest, recruiters proceeded directly with the job information session; those initially hesitant were offered a small monetary incentive (randomized between approximately USD 0.4 and USD 1.5), paid immediately after the

²⁶Initially, we held interviews on consecutive days to recruit employees for the first employment round which was planned to start in early February. Subsequently, interviews were usually held on a weekly basis. Appendix Figure B.3 reports the days on which we held interviews and the number of interviews per day.

²⁷We followed Langener et al. (2022) in designing a shortened version of Raven’s Progressive Matrices. Details about the choice experiment are provided in Appendix J.

²⁸More details are provided in Appendix B.1.

session, regardless of their subsequent interest in employment.²⁹

The job information session mirrored the previously described job interview; we collected baseline demographic information, administered the choice experiment, and provided the same terms-of-employment letter at the end. As before, the letter stated a monthly salary randomized between 55,000 NGN and 85,000 NGN (approximately USD 22 to 34), and varied in salary reliability according to the three treatment arms described in Section 4.1.

In-person recruitment was conducted in two waves. The first wave began on January 20, 2025 and ended on March 6, 2025. The second wave began on June 30, 2025 and ended on July 31, 2025. Appendix Figure B.3 shows the days on which we conducted in-person recruitment and the number of interactions per day.

Measuring Take Up. Our first primary outcome measure is individuals’ willingness to accept jobs under the conditions outlined in their respective letters detailing terms of potential employment. To credibly measure actual willingness to work, jobseekers were required to attend an *orientation day* to formally express interest in the offered positions.³⁰ Attendance at the *orientation day* provides a credible measure of job acceptance, as it constitutes a costly action clearly signaling genuine interest in employment.

4.3 Salary Delay Treatments

To identify employees’ on-the-job effort responses to delayed salary payments (target parameter 1), we randomly implemented salary delays among employees who accepted job offers in employment terms treatment arm 2. Because these job offers explicitly informed employees about the possibility of salary non-payment in any given pay cycle, we could ethically implement randomized salary delays for employees assigned to this treatment arm. The ran-

²⁹It was important that the incentive not come directly from the employing company, as this would send a strong signal about the firm. Therefore, the incentive was paid by the recruitment company.

³⁰The *orientation day* was a pre-scheduled, in-person event designed as a costly signal to elicit jobseekers’ interest. Although we provided additional job details during the session, the primary purpose was to record the jobseekers’ interest given their terms of potential employment.

domization of salary delays was stratified by age, gender, and treatment probability. Each pay cycle, employees were randomized to either receive their salaries on time or not receive payment, in which case unpaid salaries were added to the balance owed during the next pay cycle. Consequently, an employee could experience salary delays never, once, or multiple consecutive times.³¹ Multiple consecutive treatments imply salary delays lasting several pay cycles, with the maximum consecutive treatment duration restricted to three pay cycles.

While every employee ultimately received their full salary, including interest payments for any delays, the wording used in the job offers — “salaries may not be paid in a given pay cycle” — was chosen to maintain high ex-ante uncertainty. This design aimed to closely mimic salary delays at other firms, where the possibility of delays typically remains unannounced until they occur.³² By comparing the productivity of employees who have outstanding salary balances with those who receive salaries on time, we can estimate the treatment effects of salary delays.

4.4 Job Task

To accurately measure employee productivity, we designed a work task — labeling images of Lego bricks — tailored to the particular requirements of this research project.³³ This task represents a specific instance of data labeling, an increasingly common type of work in which workers manually add informational labels to individual data points (in this case, images) to identify their attributes or assign them to specific categories. Such labeled data serve as ground-truth datasets essential for training artificial intelligence models — for instance, image-recognition tools.³⁴

³¹We ensured that treatment status could not switch ‘on and off’ more than once.

³²While salary delays typically occur unexpectedly at other firms, we explicitly disclosed this possibility in our job offers for ethical reasons. See Section 4.1 for a discussion of why we believe that announcing these possibilities in the job offers does not distort our estimates.

³³We acknowledge that labeling Lego bricks is not a typical job task in this context. However, this task closely resembles skills required in many jobs: carefully following instructions, managing tasks of varying difficulty, and maintaining productivity in repetitive assignments.

³⁴We provide context about this type of work and show that it is a common task in Appendix C.4.

We digitally created 44,000 images of Lego bricks and defined six distinct categories — related to the color, type, and quantity of bricks visible in each image — according to which employees had to label the images. Categories could require one or multiple labels, resulting in a total of 10 label observations per employee for each image.³⁵ Employees performed these labeling tasks individually on provided work computers using the online data labeling platform *Labelbox*, which is designed for exactly this type of work. After completing all labels for an image, the next image appeared immediately (Appendix Figure C.3 shows an example). Importantly, all employees labeled the same dataset of images in identical order, ensuring that every image was labeled by each worker. This design eliminates variation in task content as a source of performance differences and allows us to account for learning over time when estimating treatment effects.

We designed the labeling task to include two clearly distinguishable types of images: ‘regular’ and ‘high-stakes’. Approximately 91 percent were regular images with plain white backgrounds. The remaining nine percent, designated as high-stakes, featured visually noisy, captcha-style backgrounds (an example is shown in Appendix Figure C.2), increasing their labeling difficulty. During their initial training, employees were explicitly instructed to devote additional care and effort to labeling these high-stakes images, as mistakes would be particularly costly for the company.

This pre-registered design allows us to assess effort responses on a generic task (regular images) and on a more demanding, high-stakes task where employees also have the opportunity to retaliate against the company. If outstanding salary balances induce retaliation, we would expect performance on high-stakes images to deteriorate relative to regular images, as mistakes on these images are particularly costly for the employer. By contrast, if employees seek to signal effort to their employer, we would expect performance on high-stakes images

³⁵For example, one category required employees to label the (uniform) color of the Lego bricks in the image, selecting a single label from a predefined set of colors. Another category asked employees to indicate which specific types of Lego bricks appeared in the image, requiring separate “yes” or “no” labels for each of five possible brick types. More details on the labeling categories are provided in Appendix C.2.

to be at least as high as on regular images. Further details on the image generation process and labeling categories are provided in Appendix C.

Measuring Productivity. This job task allows us to evaluate employee performance at the individual-image level, our key outcome measure. Because we digitally generated all images, we know the correct labels for each image’s categories and can directly compare employee-assigned labels to the correct ones. Our main analysis relies on two complementary, pre-registered productivity measures. The first is a continuous index measuring the total number of correct labels per image, ranging from zero (no correct labels) to ten (all labels correct). The second is a binary index equal to one if all labels on an image are correct and zero otherwise. For robustness, we also consider a productivity measure based on a pre-registered subset of the labeled images (see Appendix E.1).

4.5 Employment Arrangements

Employment took place across two distinct rounds, each designed to identify different empirical target parameters. The first employment round primarily focused on estimating employees’ on-the-job effort responses to salary delays (target parameter 1). For logistical reasons, this round began in two batches, starting on February 3 and February 10, 2025.³⁶ The second employment round, starting on August 4, 2025, was designed to assess productivity differences between worker types (target parameter 3) and further evaluate effort responses (target parameter 1) among employees with higher outside options.

Employees received instructions to resume at one of six work locations. All work locations were in proximity to each other accommodating 24 to 100 individuals. Employees were expected to work 7.5 hours per day with a 30-minute break scheduled around lunchtime. During work hours, employees only task was to label the images of Lego bricks as described in Section 4.4

³⁶Individual employees effectively started on different dates if they missed their assigned first workday or if they replaced another employee who quit shortly after starting.

On their first workday, employees received extensive training, covering computer use and detailed instruction on the labeling task, as well as logistical aspects of the job. The exact dates for the scheduled biweekly salary payments were communicated to employees. Additionally, each work location had a fixed seating arrangement, and employees were allocated to tables of four.³⁷ Communication between tables was explicitly prohibited. To enforce this rule as much as possible each table also had its own schedule, specifying start times, breaks, and closing times. These measures were implemented to address potential spillover effects of salary delays and are discussed and validated in Appendix E.2.

4.6 Supplementary Data Collections and Interventions

In addition to our primary data, we collected supplementary datasets from employee and recruitee samples, in some cases through additional interventions. These data provide further insights into employees’ experiences and responses to wage withholding, recruittees’ perceptions of job offers across the three treatment arms, and help contextualize the magnitude of the treatment effects. Specifically, we: (i) conducted weekly WhatsApp surveys to track employees’ self-reported well-being; (ii) administered an end-of-job survey to elicit employees’ beliefs about salary delays and job perceptions; (iii) implemented a bonus-payment intervention in the final week of employment round one to benchmark the magnitude of effort responses; and (iv) ran a follow-up survey with recruittees who were not hired to capture their perceptions of employment terms across treatment arms. Full details and results are reported in Appendix G.

4.7 Ethical Review

This research has undergone ethical review processes both in the United States and Nigeria, receiving IRB approval from relevant institutions. We recognize that experimentally delaying salary payments raises important ethical considerations; however, we consider the study ethically appropriate for several reasons. First, we explicitly communicated the possibility

³⁷In a few exceptions the number of employees per table was three or five.

of delayed wages clearly and transparently to all affected participants. Job offers specifically stated that salaries might not be paid in a given pay cycle, and participants made a voluntary choice whether to work under these conditions. Second, the phenomenon we study — wage withholding — is common in the studied setting as documented in this paper, meaning that participants were not exposed to risks beyond those encountered in everyday life. Third, we conducted a full debriefing with all participants; they received full information about the research and were paid any outstanding salary balance including interest payments to account for the delay period.

In addition, end-of-employment surveys indicate that employees were generally satisfied with their work experience, including those who experienced delayed salary payments. This evidence further alleviates concerns about participants’ well-being and suggests that wage withholding is not perceived as unusual in this context.

5 Data and Descriptive Statistics

5.1 Characteristics of Recruitment Samples

Over the period from January 2025 to July 2025, we recruited three distinct samples of individuals to estimate the impact of salary uncertainty on labor force participation (target parameter 2). Because we recruited for real jobs, these recruits also formed the pool of potential employees, and we later hired a subset of respondents from these samples for our two rounds of employment. First, we interviewed 638 respondents who directly applied to our posted job advertisements (Sample I).³⁸ Separately, recruiters interviewed a total of 1,079 individuals through our targeted in-person recruitment strategy.³⁹ Among these, 258 individuals agreed to participate immediately without incentive (Sample II), while 821 only agreed after receiving the monetary incentive (Sample III). Table 1 summarizes sample characteristics. From these samples, we later hired a subset for two employment rounds.

³⁸Appendix Figure B.4 shows the timing of the job-advertisement recruiting.

³⁹Appendix Figure B.3 shows the timing of the in-person recruitment; recruiters’ conversion rate from approach to interview was 44 percent.

As anticipated, the three samples differ noticeably along certain dimensions. While demographic characteristics are reasonably similar, self-reported job-search status differs substantially across groups. The share of self-employed individuals is also substantially higher among those who required a monetary incentive to participate in the job information session.⁴⁰ Attendance at the orientation day also differs noticeably across samples. While attendance was 95 percent for individuals recruited through job advertisements (Sample I), it was lower for those recruited in-person — 60 percent for the unincentivized group (Sample II) and 44 percent for the incentivized group (Sample III).

5.2 Employment Patterns and Productivity

Employment Round 1. In the first employment round, which focused on identifying employees’ effort responses to salary delays (target parameter 1), we hired 300 jobseekers from recruitee sample I (job-ad sample). The first column of Panel (a) in Table 2 shows the composition of this round; 33 employees were assigned to employment terms treatment arm 1 (control) and 267 (approximately 90 percent, as specified in our pre-analysis plan) to treatment arm 2 (salary uncertainty). Appendix Table B.9 summarizes employee demographics.

Panel (b) of Table 2 presents descriptive work patterns. On average, an employee experienced one delayed payment during the three-month spell. The absenteeism rate, unadjusted for any employees who permanently left work during the three-month period was 21 percent. Conditional on being at work, the average employee worked for 7.16 hours per day, of which 6.12 hours were spent actively labeling images. Employees differed significantly at the speed and quality with which they processed and labelled images. On average, employees labelled 15,332 images over the course of the work, with the employee in the 90th percentile labelling almost 24,000 more images than the employee at the 10th percentile. The average employee

⁴⁰We refer to this group broadly as self-employed, although it includes individuals running small-scale businesses—often not categorized as self-employment by respondents—and those engaged in casual daily labor, which respondents frequently did not perceive as employment at all. For further discussion, see Barker et al. (2024).

Table 1: Characteristics of Recruitment Samples

	Sample I	Sample II	Sample III
Characteristics			
Age	29.53	25.65	27.16
% Female	0.53	0.40	0.42
Raven’s Matrices	8.16	8.09	7.40
% Searching for a Job	0.99	0.89	0.62
% Attending Orientation	0.95	0.60	0.44
Employment Status			
Unemployed	0.51	0.47	0.26
Employed	0.10	0.11	0.15
Self-Employed	0.39	0.41	0.59
Highest Schooling			
Secondary School	0.39	0.60	0.63
Vocational Training	0.01	0.02	0.01
National Diploma	0.16	0.15	0.13
University	0.44	0.22	0.20
Observations	638	258	821

NOTE: This table reports average characteristics across our three samples of recruitees. Sample I: job-advertisement recruitees, i.e. individuals who responded to our job advertisement. Sample II: in-person recruitees who participated in the job information session without a monetary incentive. Sample III: in-person recruitees who required a monetary incentive to participate in the job information session. Appendix Table B.8 provides a more detailed version of this table, reporting additional variables, standard deviations, and percentiles.

labelled 10 percent of images correctly.

Figure 3 shows employees’ empirical learning curve. After a steep initial increase at the start of the employment spell, performance remained relatively constant over the remainder of the period. We provide summary statistics at the image-level in Appendix C.5.

Employment Round 2. The second employment round primarily focused on identifying workforce composition effects (target parameter 3). It also enabled us to estimate complementary effort responses to salary delays among employees with higher outside options (target parameter 1). We hired 300 jobseekers: 125 from recruitee sample I, 65 from recruitee


Table 2: Employee Samples

	Employment Round 1	Employment Round 2			
	Overall	Overall	Sample I	Sample II	Sample III
Panel A: Employment Sample Composition					
Arm 1: Control	0.11	0.14	0.00	0.00	0.37
Arm 2: Certainty	0.00	0.61	0.58	0.62	0.63
Arm 3: Uncertainty	0.89	0.26	0.42	0.38	0.00
Panel B: Employee Work Patterns					
Images labeled per day	255.73	207.82	205.70	204.31	212.31
Absent (0,1)	0.21	0.18	0.17	0.20	0.17
Hours at work	7.16	7.13	7.12	7.08	7.16
Hours spent working	6.12	6.15	6.16	6.10	6.17
Images labeled in total	15,332.38	8,305.79	8,220.82	8,164.52	8,485.84
Correct labels per image (0,10)	8.37	8.38	8.36	8.43	8.38
All labels correct (0,1)	0.10	0.10	0.10	0.10	0.10
Seconds spent per image	89.96	107.64	109.81	109.81	103.89
Observations	300	300	125	65	110

NOTE: This table reports employees' performance and terms of employment across the two employment rounds. For the second round, all statistics are presented separately by recruitment sample: in-person and incentivized (Sample I), in-person and unincentivized (Sample II), and job-advertisement recruits (Sample III). All employees hired in the first round were recruited through job advertisements. Panel A reports the composition of employment terms under which employees were hired. Panel B presents performance-related statistics for employees.

sample II, and 110 from recruitee sample III. Columns 2 to 5 The of Panel (a) in Table 2 present the distribution across treatment arms; 14 percent (41 employees) received treatment arm 1 (control) offers, 61 percent (182 employees) received treatment arm 2 (salary uncertainty) offers, and 26 percent (77 employees) received treatment arm 3 (salary certainty) offers.

Figure 3: Learning Curve



NOTE: This figure plots employees' empirical learning curves by employment round. The outcome is the average number of correct labels per day, smoothed using a local polynomial regression with an Epanechnikov kernel and a bandwidth of three days. The first employment round lasted roughly three months (around 60 working days), and the second lasted about two months (around 40 working days).

6 Empirical Strategy and Results

We now turn to presenting our empirical analysis and experimental findings. We begin by analyzing employees' responses at the intensive margin — that is, examining how salary delays affect the on-the-job effort and productivity of employees who continued working despite experiencing withheld wages. Subsequently, we present the extensive margin results, addressing labor force participation decisions and workforce composition effects. The empirical strategy follows our pre-analysis plan, which pre-specified the outcomes, estimation approach, and aggregation level of treatment effects.

6.1 Target Parameter 1: On-the-Job Effort Responses (Intensive Margin)

Effort Response Estimation. To estimate the effect of delayed payments on employee productivity, we exploit random treatment timing and implement the event study estimator of Borusyak et al. (2024). The estimator first fits a TWFE regression using only untreated observations and then uses the estimated parameters to predict counterfactual outcomes for treated observations. It is particularly well suited to our setting because the design of our work task allows us to condition estimation on the labelled image q , yielding granular predictions for each image.

Formally, let $G_i \in \mathcal{G} = \{1, 2, \dots, \infty\}$ denote employee i 's first treatment date t ($G_i = \infty$ if employee i is never treated). Additional, let $y_{itq}(g)$ denote the potential productivity on image q for individual i on date t when treatment started at g . We can then define event time as $l_{it} = t - g_i$.⁴¹ Event time is measured in workdays since treatment. Define the set of untreated and treated employee-day observations as $\mathcal{S}_0 = \{(i, t) : l_{it} < 0\}$ and $\mathcal{S}_1 = \{(i, t) : l_{it} \geq 0\}$, respectively. We estimate

$$y_{itq} = \alpha + \delta_q + \kappa_i + \sigma_t + \varepsilon_{itq} \quad (5)$$

via OLS on $(i, t) \in \mathcal{S}_0$ (i.e., untreated observations only). Here, y_{itq} is the outcome of interest — the continuous or binary index, described in Section 4.4 — for image q , labeled by individual i on day t . The parameters δ_q , κ_i , and σ_t denote vectors of image, individual, and date fixed effects, respectively. For treated observations $(i, t) \in \mathcal{S}_1$, we obtain predicted untreated outcomes $\hat{y}_{itq}(\infty) = \hat{\alpha} + \hat{\delta}_q + \hat{\kappa}_i + \hat{\sigma}_t$ and compute the treatment effect of salary delay at the employee-image level as $\tau_{itq} = y_{itq} - \hat{y}_{itq}$. We then aggregate these effects for each

⁴¹We can then define a binary treatment indicator $D_{it} \in \{0, 1\}$ as $D_{it} \equiv \mathbb{1}\{l \geq 0\}$.

event time g using flexible weights $\omega_{itq}^{(l)}$.⁴² Formally, we define, with $\mathcal{S}_l \equiv \{(i, t, q) : l_{it} = l\}$,

$$\tau_{g,\omega} = \sum_{(i,t,q) \in \mathcal{S}_l} \omega_{itq}^{(l)} \tau_{itq}^{(l)}. \quad (6)$$

We consider two weighting schemes in our analysis. First, each image-level observation receives the same weight, so the estimand is the average treatment effect for event-time g across all images; this is the weighting we apply throughout our main analysis. However, this weighting scheme does not account for differences in employees' labeling speed, as employees who label more images contribute proportionally more to the estimated effect. To address this, we implement a second weighting scheme that assigns equal weight to each employee-workday combination corresponding to event time g . Each image is weighted inversely proportional to the number of images labeled by that employee on that workday. We then aggregate the daily treatment effects to obtain weekly treatment effects and an overall ATE.

We compute standard errors in two ways. First, we use the conservative variance estimator proposed by Borusyak et al. (2024) to obtain standard errors for our daily event-time estimates $\tau_{g,\omega}$. We then obtain standard errors for the weekly and overall ATE aggregates by applying the delta method to our daily event-time estimates. Second, we calculate and report bootstrap standard errors, as proposed by Liu et al. (2024), for our final estimates. Since we analyze multiple productivity measures we compute sharpened q -values following Benjamini et al. (2006) to account for multiple hypothesis testing.⁴³ We report both conventional p -values and sharpened q -values in the results.

Identification of the treatment effects relies on random treatment timing, stratified by age, gender and job offer details. As a validity check, we test for pre-trends following Borusyak

⁴²We estimate treatment effects at the daily level, as specified in our pre-analysis plan, because we considered this frequency best suited to capturing heterogeneity in work performance before further aggregating treatment effects. For example, performance may be influenced by daily factors such as extreme heat, power outages, or network interruptions.

⁴³This method is suitable for positively dependent p -values, see Anderson (2008) for a discussion. This assumption is reasonable in our setting, as different productivity measures for the same image are likely correlated through underlying worker ability and effort on that particular image.

et al. (2024), testing whether all pre-treatment coefficients are equal to zero. Full details of the test are provided in Appendix D.

To ease interpretation of the estimated effects, we impose two sample restrictions in all estimations. First, we do not allow treatment status to switch “on and off.” To ensure that truthful information was conveyed to employees, individuals were randomized into treatment in each period. This implies that, in principle, an employee could be treated in one period, untreated in the next, and treated again thereafter. To prevent such switching, we restrict assignment so that individuals can only receive treatment in consecutive periods. Once an employee is randomized out of treatment after having been treated in one or more consecutive periods, they are no longer eligible to be randomized into treatment again. In addition, we exclude employees who were treated in one period but then became untreated in the following period. Since initial treatment was randomized, this exclusion effectively removes a random subsample from some weeks of the analysis. The restriction ensures, however, that we only use completely untreated individuals to estimate the fixed effects employed for imputation. Appendix XXX shows formally that this restriction does not bias the estimates. Second, we exclude the first 500 labeled images from all estimations. As shown in Figure 3, employee performance exhibits a steep initial learning curve, making early labeling decisions particularly noisy. Excluding these observations does not substantively alter the results and, if anything, leads to slightly more conservative estimates, as demonstrated in Appendix XXXX.

Estimating Absenteeism and Day-Specific Outcomes. In addition to productivity measured at the image level, we also estimate absenteeism and other outcomes defined at the workday level, such as total time spent labeling and total work time. Our estimation strategy follows the approach outlined above, again using the imputation estimator of Borusyak et al. (2024), with the difference that outcomes are defined at the day rather than the image

level.⁴⁴

Effort Results. We find that salary delay increases effort in image classification by about 0.5 percent relative to the control mean, with the dynamic patterns closely resembling the simulations from our theoretical framework. Figure 4 presents weekly event-study estimates for both the continuous (Panel (a)) and binary indices (Panel (b)), obtained by aggregating the daily treatment effects from equation (6) to the weekly level.⁴⁵ Over the course of the experiment, employees increased their effort in response to longer salary delays. The pre-treatment trends in both panels are flat and close to zero. For the continuous index, there is an increase in the final week before treatment, but the coefficient is statistically insignificant at conventional confidence levels.

Table 3 aggregates the daily treatment effects from equation (6) into a single ATE. Columns (1) and (2) present results for the continuous and binary indices, respectively, for the full sample of images. Columns (3) and (4) are analogous but restrict the sample to the ‘high-stakes’ images, for which employees were encouraged to exert extra effort.⁴⁶ The treatment effects are very similar across the two image samples, indicating that employees who experienced salary delay maintained a consistently higher level of effort than those who received their salary on time.

The coefficient on the continuous index quantifies the change in correct labels per image, while the coefficient on the binary index quantifies the change in the probability of flawless labeling. For example, the coefficient in column (1) indicates that employees’ performance increases by 0.045 correct labels per image, on average, during periods of payment delay. Relative to a baseline mean of 8.42 correct labels, this corresponds to a productivity increase

⁴⁴Specifically, we estimate

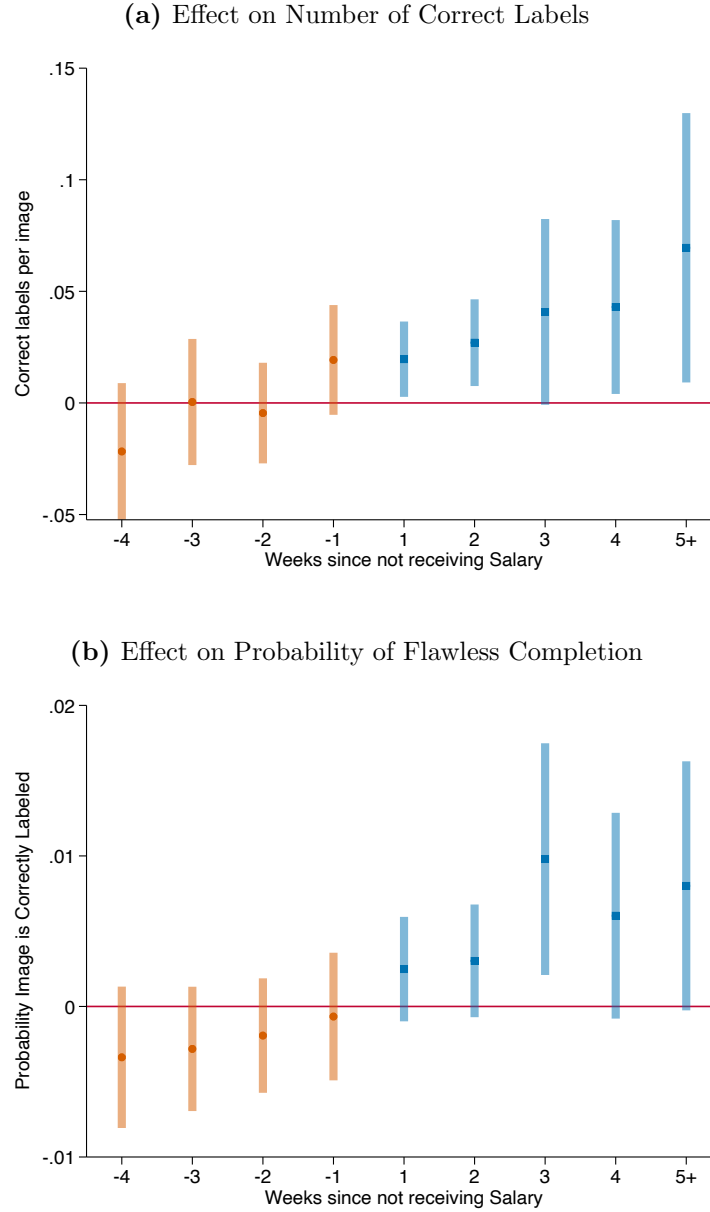
$$y_{it} = \alpha + \kappa_i + \sigma_t + \varepsilon_{it} \quad (7)$$

via OLS on $(i, t) \in \mathcal{S}_0$. Here, κ_i and σ_t are defined as before, and y_{it} denotes weekday-specific outcomes.

⁴⁵As pre-registered, we considered this the most appropriate level ex-ante for analyzing dynamic productivity responses to salary delay.

⁴⁶See Section 4.4 for a more detailed discussion.

Figure 4: Event Study Graphs: Estimates from Equation (6)



NOTE: This figure plots the dynamic effect of being owed salary payments on employees' effort measures. Panel (a) presents results for the number of correct labels per image, while Panel (b) shows results for the binary indicator of whether an image was labeled entirely correctly. The blue squares show the estimates of equation (6) aggregated at the weekly level, standard errors are calculated using the conservative variance estimator proposed in Borusyak et al. (2024). The orange circles show estimates from a separate regression testing for pre-trends as suggested by Borusyak et al. (2024) (these are not from the same estimation and coefficients are not relative to a single omitted time period as is often the case in figures like this one), see Appendix D for details. Standard errors are initially clustered at the individual level and aggregated using the delta method.

of 0.53 percent. This, in turn, corresponds to an increase in the probability of flawless labeling of 0.006 percentage points, as shown in column (2). Relative to an 11 percent baseline probability of flawless labeling, this implies a treatment effect of 5.4 percent.

Table 3: Treatment Effects of Wage Withholding on Worker Effort

	All Images		High-Stakes Images	
	Continuous Index	Binary Index	Continuous Index	Binary Index
ATE	0.045 (0.016)*** [0.018]***	0.006 (0.003)** [0.003]**	0.048 (0.016)*** [0.019]***	0.008 (0.003)*** [0.003]***
Observations (First Stage)	2,166,877	2,166,877	196,994	196,994
Observations (Imputed)	675,184	675,184	61,209	61,209
Individuals	297	297	297	297
SE Cluster	Individual	Individual	Individual	Individual
Mean of Dep. Var.	8.43	.11	8.41	.11
Q-Value	0.010	0.016	0.011	0.011

NOTE: This table reports the effect of being owed salary payments on employees' effort measures from equation (6), aggregated into a single ATE. Standard errors based on the conservative variance estimator of Borusyak et al. (2024), clustered at the individual level and aggregated using the delta method, are shown in parentheses. Bootstrap standard errors with 500 replications, following Liu et al. (2024), are reported in brackets. 'First-stage observations' refer to the number of observations used to estimate equation (5), while 'imputed observations' indicate the number of counterfactual outcomes generated to estimate treatment effects. We report sharpened q -values accounting for multiple hypothesis testing across our four productivity measures, based on the p -values in parentheses.

To contextualize these magnitudes, we compare them to a one-standard-deviation change in performance for each index. Across employees, the standard deviation of performance is 0.39 for the continuous index and 0.04 for the binary index (Panel C of Table 2). These estimates imply that, in the initial weeks of experiencing salary delay, employees improve performance by approximately 11.5 to 15.5 percent of a standard deviation, depending on the productivity measure considered.⁴⁷

⁴⁷The experiment cannot speak to any results beyond the experimental six-week window.

As a second benchmark, we compare the estimated performance increase to the increase induced by a bonus payment. In a separate intervention, we offered a bonus of ten NGN per correctly labeled image, which increased performance by 0.014 percentage points. Under the assumption of a linear relationship between bonus amount and performance, the estimated treatment effect of salary delay is equivalent to offering a bonus of approximately five NGN per correct image.⁴⁸

The Appendix contains multiple versions of Figure 4 and Table 3 incorporating a range of robustness checks and complementary results. First, Appendix E.2 addresses potential spillover effects of salary delay across employees — arising from the concern that individuals with different treatment status may share the same work location — and shows that the results are not sensitive to accounting for such spillovers. Second, we show that the results are robust to using an alternative productivity index based only on a pre-registered subset of the labels in Appendix E.1. Third, Appendix E.3 reports results using the alternative weighting scheme for calculating daily treatment effects $\tau_{g,\omega}$, outlined earlier in this section, and shows that results are largely robust to the choice of weighting. Fourth, we find that the results are not meaningfully affected by the level of treatment penetration within a work location, as reported in Appendix E.4. Fifth, we show that treatment effects are larger in magnitude when the sample is restricted to the most difficult images — those with the lowest average labeling performance in the untreated sample. Given high overall labeling performance, additional effort should matter more for the most difficult images, and Appendix E.5 provides evidence consistent with this expectation.

Heterogeneity of Effort Responses. To understand what drives the effort responses to salary delay and if firms may face retaliation from a subset of employees, we investigate heterogeneity of treatment effects. Specifically, we estimate versions of Equation (6) in which estimands are separately aggregated for binary subgroups. The results are presented

⁴⁸Five NGN corresponds to around 1.2 percent of the hourly minimum wage.


in Figure 5. We begin by analyzing subgroups that may differ in their outside options and, therefore, in how costly it is to forgo outstanding salaries or lose their jobs. In particular, we examine treatment effects by pre-treatment performance, postsecondary education, and gender.

Employees' pre-treatment performance is informative about their outside options. Lower-performing workers are less productive and therefore less attractive to alternative employers, limiting their outside options. Similarly, we interpret postsecondary education as a proxy for higher outside options, as better-educated workers are more likely to find alternative employment. In addition, men may have better outside opportunities than women in this labor market, since much of casual employment involves physically demanding work that limits women's access to these jobs. Following the logic of our theoretical framework, groups with lower outside options should increase their effort more strongly in response to withheld wages. These workers are more dependent on their current job, and forgoing outstanding salary balances is more costly for them. They therefore have stronger incentives to intensify effort when salary payments are delayed. At the same time, workers with higher outside options could be considered more likely to retaliate over unpaid wages.


We find evidence consistent with the prediction that effort increases should be particularly pronounced among workers with low outside options while not finding any evidence of retaliation among workers with higher outside options. First, workers who performed below the median before treatment increased their effort significantly more while wages were outstanding than those who performed above the median ($p < 0.01$). Panel (a) shows the dynamic treatment effects: labeling performance increases significantly for below-median performers, while we find a precise null effect for those above the median. Second, splitting the sample by postsecondary education yields similar results. Workers without postsecondary education show a significantly stronger effort response than those with postsecondary education ($p < 0.05$). Dynamic results are presented in Panel (b). Third, we find suggestive evidence

Figure 5: Treatment Effect Heterogeneity


(a) Heterogeneity by Pre-Treatment Performance




(b) Heterogeneity by Schooling



(c) Heterogeneity by Gender



(d) Heterogeneity by WTP for Salary Certainty



NOTE: This figure illustrates treatment effect heterogeneity. Panel (a) presents heterogeneity by pre-treatment performance, splitting the sample into above- and below-median performers. Panel (b) shows heterogeneity by post-secondary schooling status. Panel (c) reports heterogeneity by gender, and Panel (d) illustrates heterogeneity by WTP for higher salary certainty.

that women increase their effort slightly more than men in response to wage withholding, but this difference is not significant. Panel (c) displays the dynamic results.

Next, we examine whether heterogeneity in workers' valuation of salary certainty matters for their responses to unpaid wages. To do so, we draw on the discrete choice experiment conducted during the job interviews and information sessions. In the choice experiment, job-seekers made incentive-compatible choices between hypothetical jobs, allowing us to estimate

a mixed logit model, and derive individual-level WTP values for higher salary certainty.⁴⁹ We may think that different valuations for higher salary certainty translates into heterogeneous responses when wages are delayed. However, we find little evidence in favor of this hypothesis. Panel (d) plots the dynamic responses for below- and above-median valuation subgroups. While workers with higher WTP for salary certainty exhibit a slightly more pronounced effort response, we cannot reject equality of responses on average. Importantly, neither group retaliates against the employer.⁵⁰

Absenteeism and Hours Worked Results. Table 4 shows estimation results for absenteeism and related outcomes. All standard errors are clustered by individual employee and combined using the delta method. Column (1) shows the estimation results for the effect of salary delay on being absent for an entire workday. The coefficient of 0.026 implies that employees were 2.6 percentage points more likely to be absent during the time (two weeks) of the pay cycle in which they did not receive their salary for the first time. This effect, however, is not significant at a reasonable confidence level.

Column (2) shows the effect for total work hours, conditional on showing up to work. Employees have fixed start and end time each day, which correspond to a workday of 7.5 hours. Employees could, however, resume late or leave early reducing the effective number of work hours. Some employees do so, as the average number of work hours is 7.19. We can infer this time from the first and last image that employees label. Column (2) shows that there is no significant effect on total work hours caused by salary delay, conditional on showing up to work.

Column (3) shows the effect for time spent labelling. We can infer the time employees spend on each image, as *Labelbox* tracks this time. Importantly, the timer pauses after five minutes

⁴⁹Details on the methodology and setup of the choice experiment, and the distribution of individual level WTP values are provided in Appendix J.

⁵⁰For all considered outcomes, Appendix Table E.16 reports subgroup ATEs and equality tests, and Appendix Figure E.6 replicates the event-study results using the binary productivity index.

of inactivity. Hence, if an employee is not working on their job task for a while, the timer will stop. We sum the labelling times for every workday and obtain the time employees spent labelling images for a given day. Our results suggest that there is no meaningful effect of salary delay on time spent labelling.

Table 4: Effect of Salary Delay on Employees’ Absenteeism and Effort

	Absent for Workday	Total Work Hours	Time Spent Labelling
ATE (One Pay Cycle Delay)	0.026 (0.020)	-0.008 (0.033)	-0.012 (0.048)
ATE (Two Pay Cycles Delay)	0.018 (0.031)	0.010 (0.053)	-0.064 (0.083)
ATE (Three Pay Cycles Delay)	-0.033 (0.048)	-0.045 (0.109)	-0.075 (0.156)
Observations	14,289	12,073	12,073
Individuals	300	300	300
SE Cluster	Individual	Individual	Individual
Mean of Dep. Var.	.181	7.19	6.17

NOTE: This table shows the effect of being owed salary payments on worker absenteeism and effort measures. The treatment effects are estimated at the worker-workday level using the Borusyak et al. (2024) imputation estimator. The daily treatment effects are combined into a single parameter spanning one pay cycle using the delta method. Standard errors are initially clustered at the individual level. This table shows the effect of salary delay on absenteeism in column (1), on total work hours in column (2) and on the total time spend labelling during the work hours in column (3).

6.2 Target Parameter 2: Labor Force Participation (Extensive Margin)

Empirical Strategy. Our goal is to identify and estimate the ATE of offering a credible salary guarantee on individuals’ willingness to accept wage employment. This parameter, θ is of both theoretical and policy relevance. It is theoretically relevant because one of the key predictions of our model is that $\theta \geq 0$. Specifically, $\theta > 0$ for individuals with sufficiently high outside options who stay out of the wage sector given the present uncertainty of salary payments, and $\theta = 0$ for individuals who would accept employment under current conditions

in Nigeria. From a policy perspective, θ is important because it quantifies the proportion of the workforce that remains unwilling to take up jobs in the current economic environment but would do so if salaries were more secure.

To estimate this effect, we employ the following specification:

$$D_i = \alpha + \lambda Z_{i2} + \theta Z_{i3} + \nu \text{wage}_i + \sigma_t + \varepsilon_i. \quad (8)$$

Here, D_i is a binary indicator equal to one if individual i attended the orientation day and zero otherwise. $Z_{iz} = \mathbb{1}\{Z_i = z\}$ is an indicator equal to one if individual i was assigned to employment terms treatment arm $z \in \{1, 2, 3\}$, with Treatment Arm 1 (Control Arm) serving as the omitted reference category. We also include a continuous measure of the offered wage (wage_i); its coefficient ν captures the direct effect of salary amount on an individual's decision to attend the orientation day. We add date fixed effects σ_t for precision.⁵¹

By construction of our RCT, the joint assignment of employment terms and wage is independent of potential attendance $D_i(z, w)$. This ensures that the following identifying assumption holds by design, so that θ , λ and ν are identified:

Assumption E1.

$$(Z_i, \text{wage}_i) \perp\!\!\!\perp \{D_i(z, w) : z \in \{1, 2, 3\}, w \in \mathcal{W}\}.$$

We estimate equation (8) using ordinary least squares (OLS) and a logistic regression via maximum likelihood estimation (MLE). Our primary coefficient of interest, θ , captures the causal effect of receiving a job offer corresponding to Treatment Arm 3 (Salary Certainty Arm), relative to the control group (Employment Terms Treatment Arm 1). Additionally, we use the ratio of coefficients $\frac{\theta}{\nu}$ to derive a WTP measure. Specifically, this ratio expresses the impact of receiving the high-certainty job offer relative to the impact of a salary increase, thus

⁵¹Appendix H reports results without date fixed effects

quantifying how large a salary increase would need to be to generate an equivalent increase in orientation-day attendance as the high-certainty employment terms. We interpret the coefficient ratio $\frac{\theta}{\nu}$ as a monetary estimate of participants' WTP for increased salary certainty.

We additionally validate our WTP estimates using the discrete-choice (conjoint) experiment conducted during the initial interview and job information session. In this experiment, participants made incentive-compatible choices among employment descriptions that systematically varied in terms of salary amount and salary certainty. Detailed methodological information are provided in Appendix J.

Results. The treatment effect estimates of θ are in line with our theoretical predictions: receiving a salary guarantee increases willingness to accept the job, as measured by orientation-day attendance. Importantly, this effect is driven by individuals who were only willing to attend the job information session when offered a monetary incentive. Table 5 reports the effect of salary certainty on willingness to accept jobs among individuals who do not self-select into employment — those who were recruited in person rather than responding to the job advertisement. Column (1) and (2) show the effect of receiving a job offers corresponding to treatment arm 2 (salary certainty arm) and treatment arm 3 (uncertainty arm) relative to the control offer. Column (3) uses only the sub-sample that received an incentive to participate in the job information session while column (4) uses only the sample that participated in the job information session without incentive.

The coefficients on *Salary Guarantee* correspond to our empirical estimate of θ . For instance, the coefficient in column (2) implies that receiving a letter outlining potential terms of employment with a salary guarantee through the automated payment system increased orientation-day attendance by 11.1 percentage points. This is a substantial effect: participation in the untreated in-person recruited sample was 44.7 percent, so the treatment effect corresponds to an effect size of approximately 25 percent.

Columns (3) and (4) provide evidence that the effect of salary certainty operates primarily

through individuals who required a monetary incentive to attend the job information session. For this subsample, the coefficient on the salary guarantee remains nearly unchanged (11.8) and highly significant, with treatment effects becoming larger in magnitude given the lower baseline attendance of 40.5 percent.⁵² By contrast, the effect loses statistical significance among individuals who were willing to participate in the job information session without an incentive.⁵³ This pattern suggests that some individuals are discouraged from wage employment by uncertainty around salary payments but become willing to engage in it once this uncertainty is reduced. At the same time, salary guarantees appear to have little impact on those already prepared to work under current economic conditions.

Table 5: Effect of Salary Guarantee on Job Take-Up

	Logit AME	LPM	LPM	LPM
Salary Guarantee (ATE)	0.121*** (0.033)	0.111*** (0.032)	0.118*** (0.037)	0.078 (0.068)
Salary Uncertainty (ATE)	0.005 (0.038)	0.002 (0.037)	-0.011 (0.043)	-0.015 (0.073)
Salary (1,000 NGN)	0.005*** (0.002)	0.005*** (0.001)	0.006*** (0.002)	-0.001 (0.003)
Observations	1,079	1,079	821	258
R-Square		0.29	0.27	0.49
Date FE	Yes	Yes	Yes	Yes
Mean of Dep. Var.	.447	.447	.405	.596
Sample	All	All	Incentivised	No Incentive

NOTE: This table reports estimates of equation (8). The results compare receiving terms of potential employment under treatment arm 2 (salary certainty) and treatment arm 3 (salary uncertainty) relative to the control arm (arm 1), which is the omitted category in all specifications. Standard errors are heteroskedasticity-robust. The sample is restricted to in-person recruits. Column (1) presents average marginal effects from a logit estimation, while columns (2) to (4) report results from a linear probability model.


Additionally, we compare the effect of salary certainty with that of higher wages on orientation-day attendance. The coefficient on *Salary* in column (2) shows that a 1,000 NGN increase

⁵²The lower baseline attendance in this subsample also indicates that the treatment effects are not merely driven by the provision of a monetary incentive.

⁵³Baseline participation in this subsample was higher at 59.6 percent, consistent with a greater willingness to work under current circumstances.

in monthly wages raises participation by 0.05 percentage points. Based on this estimate, we compute the ratio of the coefficients on *Salary Guarantee* and *Salary* to translate the effect of salary certainty into an equivalent wage increase. The first two bars in Figure 6 show these ratios, corresponding to the estimates in columns (1) and (2) of Table 5.

Figure 6: WTP for higher Salary Certainty



NOTE: This figure provides estimates of the willingness to pay for a salary guarantee. The first two bars show the ratios of the coefficients on Salary Guarantee and Salary based on the estimates in columns (1) and (2) of Table 5. The third bar reports the corresponding ratio from a conditional logit estimation (equation (21), Appendix J). Standard errors are calculated using the delta method in all cases. The conditional-logit estimate from the choice experiment (bar 3) is statistically indistinguishable from the first two estimates ($p = 0.64$ and $p = 0.74$ respectively).

We obtain an estimate of jobseekers' willingness to pay of approximately USD 15 (about 22,500 NGN). The confidence intervals around the estimates are wide, however, so we cannot rule out a broader range of values. As a complementary source of evidence, we draw on the discrete choice experiment conducted during the job interviews and information sessions. In the choice experiment, jobseekers made incentive-compatible choices between hy-

pothetical jobs that varied in salary certainty and wage levels, corresponding to the three treatment arms and analogous to the actual terms of employment. Estimating a standard conditional logit model, we find a willingness-to-pay of about USD 17.50, remarkably close to the orientation-day attendance estimate but with much tighter confidence intervals.⁵⁴ These two approaches provide strong evidence that individuals place substantial value on salary certainty, with an implied willingness-to-pay between USD 15 and 18 (22,500-26,500 NGN). This corresponds to more than 100 percent of the weekly median wage and about one-third of the legal monthly minimum wage, which is widely regarded as a good salary.

6.3 Target Parameter 3: Workforce Composition

Having established that salary uncertainty influences labor-force selection, we now assess whether it also affects workforce composition (Target Parameter 3). Specifically, we examine whether it alters the types of workers willing to accept jobs and whether this, in turn, has productivity implications for firms. To address these questions, we use two complementary strategies. First, we characterize the individuals induced to accept our job offer by the salary guarantee—the compliers in a standard LATE framework (Imbens and Angrist, 1994). Our main analysis then uses data from the second employment round to compare the productivity of in-person recruits hired under the guarantee with that of job-advertisement applicants who self-selected into the position.

Estimation of Complier Characteristics. As outlined in Section 6.2, offering a salary guarantee increases job take-up among the in-person recruited sample. This raises the question of how salary guarantees affect workforce composition by inducing additional individuals at the extensive margin to accept wage employment. Following the terminology of the standard LATE framework (Imbens and Angrist, 1994), we distinguish three groups: always takers, who accept employment regardless of the guarantee; compliers, who accept only if the salary guarantee is offered; and never takers, who decline employment even with the

⁵⁴Details on the methodology and exact setup of the choice experiment are provided in Appendix J.

salary guarantee. Relying on Abadie (2002, 2003) we can characterize the compliers using an instrumental variable framework where we use a slightly modified version of Equation (8) as the first stage. Specifically we estimate the following 2SLS framework:

$$c_i \times D_{id} = \gamma_0 + \gamma_d D_{id} + \delta_{w,t} + u_i \quad (9)$$

$$D_{id} = \pi_0 + \pi_1 Z_{i3} + \delta_{w,t} + e_i. \quad (10)$$

In a slight extension of notation, $D_{id} = \mathbb{1}\{D_i = d\}$ is an indicator equal to one if individual i 's attendance at the orientation corresponds to $d \in \{0, 1\}$, with $d = 1$ denoting attendance and $d = 0$ non-attendance. As before, Z_{i3} denotes assignment to treatment arm 3. To improve precision and strengthen the first stage, we nonparametrically control for date t , wage level w , and their interactions, thereby saturating the controls which is required for a LATE interpretation of 2SLS with controls (Blandhol et al., 2022). Accordingly, δ_{wt} denotes the full set of wage-by-date interaction dummies ($w \in \mathcal{W}$, $t \in \mathcal{T}$), i.e., the saturated version of the wage and σ_t controls used in (8). Results from a specification without controls are reported in Appendix I. Let c_i denote a characteristic of individual i (Raven's score, age, schooling, or job-search status). We interact c_i with the binary attendance indicator D_{id} and regress this constructed outcome on orientation attendance, instrumented by assignment to employment terms treatment arm 3.

Under standard LATE assumptions — random assignment, exclusion, monotonicity, and relevance (discussed in detail in Appendix I) — the coefficient γ_d identifies the mean level of characteristic c among compliers with attendance status d . This allows us to characterize compliers who attended as well as those who did not attend the orientation day.⁵⁵ We pool the two complier groups and report their average characteristics, with standard errors computed using the delta method. Related applications of this approach appear, for example,

⁵⁵Compliers who did not attend the orientation day are those who would have attended had they received a terms-of-employment offer with a salary guarantee. See Appendix I for details on the identification of this group.

in Autor and Houseman (2005) and Angrist et al. (2023), and additional details are provided in Appendix I. We use heteroskedasticity-robust standard errors.

Estimating Productivity Comparisons. Our second strategy examines how salary uncertainty affects workforce composition using evidence from the second employment round. In this round, both job-ad and in-person recruitees were offered contracts that included a salary guarantee. This allows us to compare individuals recruited in person — where we previously observed large selection effects from the guarantee — with jobseekers who responded to job advertisements, a group characterized by an extremely high take-up rates regardless of contract terms. To carefully assess any performance differences between the two groups that might have implications for firms’ productivity, we estimate a dynamic specification comparing a range of daily productivity measures. Specifically, we estimate the following regression:

$$y_{itq} = \alpha + \gamma R_i + \sigma_t + \gamma_t (R_i \times \sigma_t) + \varepsilon_{itq}. \quad (11)$$

Here, y_{itq} denotes the outcome of interest — the continuous or binary index described in Section 4.4 — for image q , labeled by individual i on day t . Additional outcomes — absenteeism, number of correctly labelled images, work hours, and effective work time — are defined at the daily level, in which case the dependent variable becomes y_{it} . The vector σ_t represents a date fixed effects and R_i is an indicator for in-person recruitment. The coefficients γ_t trace how performance differences between the two groups evolve over time. We restrict the analysis to the first month of employment — before any individuals were exposed to salary delays — and to employees whose job offers did not specify potential salary nonpayment.

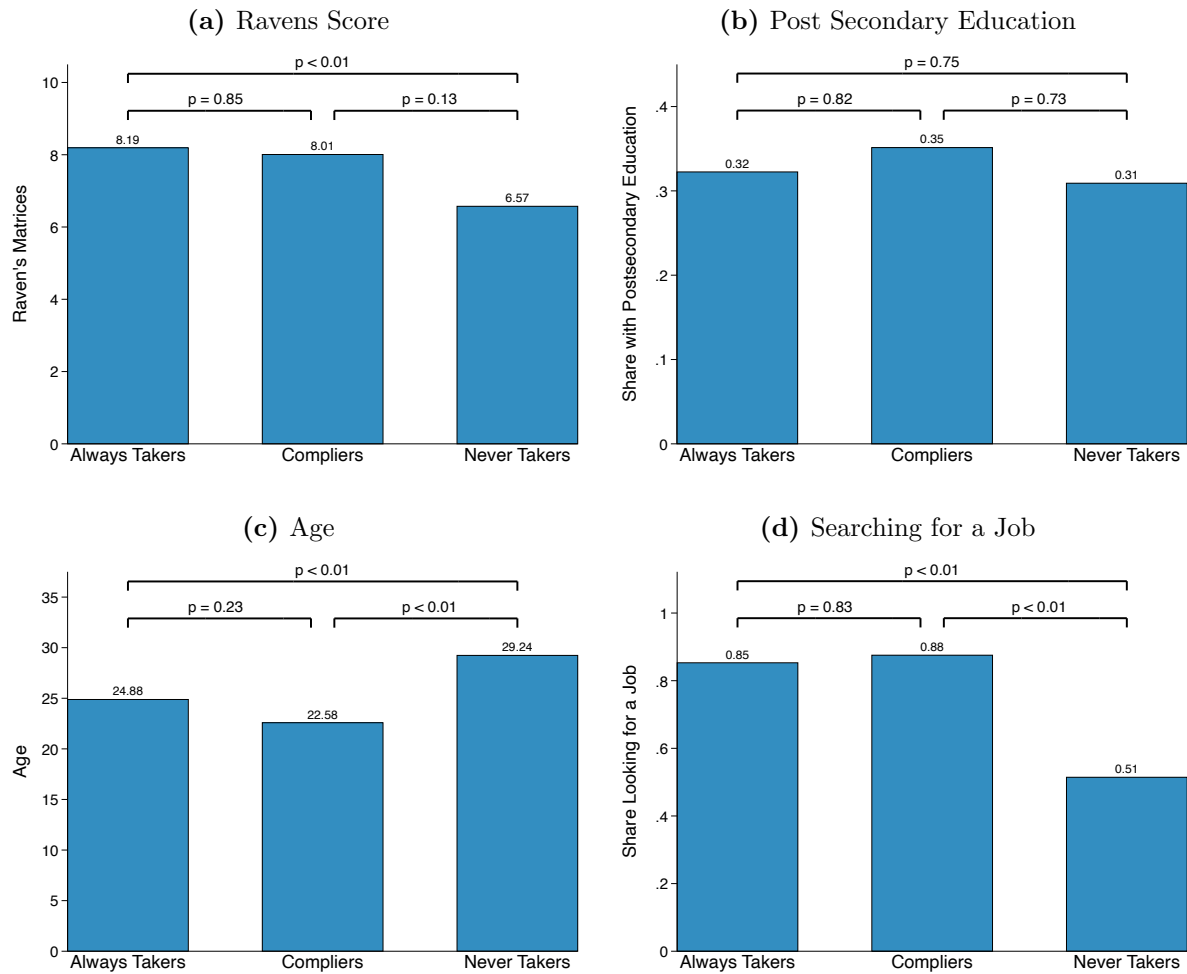
Complier Characteristics Results. Compliers are much more similar to always-takers than to never-takers, as shown in Figure 7. We report results for four key characteristics: Raven’s score, attainment of post-secondary education, age, and jobsearch status.

Across all four characteristics, compliers are statistically indistinguishable from them, in contrast to never-takers, who differ noticeably across several dimensions. Panel (a) shows that compliers' Raven's scores are nearly identical to those of always-takers and somewhat higher than those of never-takers, though the latter difference is not statistically significant. In Panel (b), post-secondary education attainment is similar across all three groups, with no significant differences. Panel (c) indicates that compliers are younger on average than both always- and never-takers, and the difference relative to never-takers is statistically significant. Finally, Panel (d) reveals a stark contrast in job-search status: compliers, like always-takers, are highly likely to report active job search, whereas never-takers are much less likely.

Productivity Comparison Results. We find no meaningful performance difference between the in-person and job-advertisement recruitees. Figure 8 presents results from estimating equation (11) for all six outcome measures, plotting daily averages for both groups. Panels (a) and 8b show estimates of equation (11) at the employee-image level using the continuous and binary productivity measures as outcomes. Productivity does not differ significantly on any given day, and overall averages are statistically indistinguishable.


Panels (c)-8f report estimates from equation (11) at the employee-day level. Again, both groups perform very similarly across all dimensions. If anything, in-person recruitees are slightly more likely to be absent on average ($p < 0.05$).

Figure 7: Complier Characteristics



NOTE: This figure shows characteristics of always takers, compliers, and never takers.

Figure 8: Workforce Composition Effects



NOTE: This figure compares performance of in-person recruits with job-ad recruits. Panel (a) shows the number of correct labels per image, and Panel (b) displays the probability of labelling an image correctly. Both panels are estimated at the employee-image level. Panels (c), (d), (e), and (f) are estimated at the employee-workday level and report daily absenteeism, number of correctly labelled images, total hours worked, and effective hours worked, i.e. time spent labelling.

7 Conclusion

This paper highlights the prevalence and consequences of overdue and unpaid salary payments, a common but underexplored feature of labor markets in low-income countries. Using original data from Nigeria, we first document that 30 percent of workers report experiencing payment difficulties with an employer. We then offer a theoretical explanation for why wage arrears can emerge, particularly when contractual enforcement is weak. Moving beyond the survey evidence we conduct a field experiment in Lagos that allows us to closely track employees' productivity and jobseeker responses to salary uncertainty. The results show that, conditional on attendance, workers exert a small but statistically significant increase in effort when salaries are delayed, while absenteeism and hours worked remain unaffected. Finally, individuals initially reluctant to pursue wage employment are substantially more likely to accept job offers that credibly guarantee salary certainty. Together, these findings underscore how weak enforcement of wage payments shapes both firm and worker behavior in low-income settings.

References

- Abadie, Alberto**, “Bootstrap Tests for Distributional Treatment Effects in Instrumental Variable Models,” *Journal of the American Statistical Association*, 2002, *97* (457), 284–292.
- , “Semiparametric instrumental variable estimation of treatment response models,” *Journal of Econometrics*, 2003, *113* (2), 231–263.
- Akcigit, Ufuk, Harun Alp, and Michael Peters**, “Lack of Selection and Limits to Delegation: Firm Dynamics in Developing Countries,” *American Economic Review*, 2021, *111* (1), 231–275.
- Akomolafe, Mohammed Akinola**, “The Rule of Law and its Practice in Nigeria: An Assessment,” *Makurdi Owl Journal of Philosophy (MAJOP)*, 2021, *2* (2), 13–28.
- Akuopha, Ochuko**, “Unpaid Salaries: Pipeline Surveillance Shut Down Oil Facilities in Delta,” April 2023. Vanguard (Ughelli), accessed October 17, 2025.
- Alfandari, Gilles and Mark Edwin Schaffer**, ““Arrears” in the Russian Enterprise Sector,” in Simon Commander, Qimiao Fan, and Mark Edwin Schaffer, eds., *Enterprise Restructuring and Economic Policy in Russia*, Washington, DC: World Bank, 1996, pp. 87–139.
- Anderson, Michael L.**, “Multiple Inference and Gender Differences in the Effects of Early Intervention: A Reevaluation of the Abecedarian, Perry Preschool, and Early Training Projects,” *Journal of the American Statistical Association*, 2008, *103* (484), 1481–1495.
- Angrist, Joshua, Peter Hull, and Christopher Walters**, “Chapter 1 - Methods for measuring school effectiveness,” in Eric A. Hanushek, Stephen Machin, and Ludger Woessmann, eds., *Handbook of the Economics of Education*, Vol. 7, Elsevier, 2023, pp. 1–60.
- Ashraf, Nava, Oriana Bandiera, and Katharina Jack**, “No Margin, No Mission? A Field Experiment on Incentive Pay for Teachers in Zambia,” *Journal of Public Economics*, 2014, *120*, 60–71.
- , – , **Sanghamitra Lee, and Emily Nix**, “Togetherness to Win: Experimental Evidence on Team Incentives in Zambia,” *American Economic Journal: Applied Economics*, 2018, *10* (2), 1–35.
- Autor, David H. and Susan N. Houseman**, “Do Temporary Help Jobs Improve Labor Market Outcomes for Low-Skilled Workers? Evidence from Random Assignments,” Staff Working Paper 05-124, W.E. Upjohn Institute for Employment Research October 2005. Revised from January 2005.
- Bandiera, Oriana, Iwan Barankay, and Imran Rasul**, “Incentives for Managers and Inequality among Workers: Evidence from a Firm-level Experiment,” *Quarterly Journal of Economics*, 2007, *122* (2), 729–773.

- Barker, Nathan, Inbar Amit, Alison Andrew, Robert Garlick, Kate Orkin, and Carol Nekesa**, “Flexibility versus Performance: The Determinants of Labor Contracts in Nairobi, Kenya,” November 2024. Job Market Paper.
- Benjamini, Yoav, Abba M. Krieger, and Daniel Yekutieli**, “Adaptive linear step-up procedures that control the false discovery rate,” *Biometrika*, 09 2006, *93* (3), 491–507.
- Bernhardt, Annette, Ruth Milkman, Nik Theodore, Douglas Heckathorn, Mirabai Auer, James DeFilippis, Ana Luz González, Victor Narro, Jason Perelshteyn, Diana Polson, and Michael Spiller**, “Broken Laws, Unprotected Workers: Violations of Employment and Labor Laws in America’s Cities,” Technical Report, National Employment Law Project, UCLA Institute for Research on Labor and Employment 2009.
- Bewley, Truman F**, “Why not cut pay?,” *European Economic Review*, 1998, *42* (3), 459–490.
- Blandhol, Christine, John Bonney, Magne Mogstad, and Alexander Torgovitsky**, “When is TSLS Actually LATE?,” Working Paper 29709, National Bureau of Economic Research January 2022.
- Blattman, Christopher and Stefan Dercon**, “The Impacts of Industrial and Entrepreneurial Work on Income and Health: Experimental Evidence from Ethiopia,” *American Economic Journal: Applied Economics*, 2018, *10*, 1–38.
- Bloom, Nicholas and John Van Reenen**, “Why Do Management Practices Differ across Firms and Countries?,” *Journal of Economic Perspectives*, 2010, *24*, 203–224.
- , **Benn Eifert, Aprajit Mahajan, David McKenzie, and John Roberts**, “Does Management Matter? Evidence from India,” *The Quarterly Journal of Economics*, 2013, *128*, 1–51.
- Boehm, Johannes and Ezra Oberfield**, “Misallocation in the Market for Inputs: Enforcement and the Organization of Production,” *The Quarterly Journal of Economics*, 2020, *135*, 2007–2058.
- Borusyak, Kirill, Xavier Jaravel, and Jann Spiess**, “Revisiting Event-Study Designs: Robust and Efficient Estimation,” *The Review of Economic Studies*, 02 2024, *91* (6), 3253–3285.
- Boyarchuk, Dmytro, Lilia Maliar, and Serguei Maliar**, “The consumption and welfare implications of wage arrears in transition economies,” *Journal of Comparative Economics*, 2005, *33* (3), 540–564.
- Breza, Emily and Supreet Kaur**, “Labor Markets in Developing Countries,” Working Paper 33908, National Bureau of Economic Research June 2025.
- , – , and **Yogita Shamdasani**, “The Morale Effects of Pay Inequality,” *Quarterly Journal of Economics*, 2018, *133* (2), 611–663.

- Buehren, Niklas, Virginia Ceretti, Ervin Dervisevic, Markus Goldstein, Leora Klapper, Tricia Koroknay-Palicz, and Simone Schaner**, “Salary Delays and Overdrafts in Rural Ghana,” *AEA Papers and Proceedings*, May 2018, 108, 449–52.
- Carranza, Eliana, Alexia Donald, Florentin Grosset, and Supreet Kaur**, “The Social Tax: Redistributive Pressure and Labor Supply,” *NBER Working Paper*, 2022.
- Cassan, Guilhem, Daniel Keniston, and Tatjana Kleineberg**, “A Division of Laborers: Identity and Efficiency in India,” 2022. Working Paper.
- Cefala, Lauren, Supreet Kaur, Heather Schofield, and Yogita Shamdasani**, “Habit Formation in Labor Supply,” *Working Paper*, 2024.
- Chemin, Matthieu**, “Does Court Speed Shape Economic Activity? Evidence from a Court Reform in India,” *The Journal of Law, Economics, and Organization*, 2012, 28, 460–485.
- Clarke, Simon**, “Trade unions and the non-payment of wages in Russia,” *International Journal of Manpower*, 1998, 19 (1/2), 68–94.
- Clemens, Jeffrey and Michael R. Strain**, “Understanding “Wage Theft”: Evasion and avoidance responses to minimum wage increases,” *Labour Economics*, 2022, 79, 102285.
- Diamond, Jack and Christian Schiller**, “Expenditure Arrears,” in Mario I. Blejer and Adrienne Cheasty, eds., *How to Measure the Fiscal Deficit*, Washington, DC: International Monetary Fund, 1993, pp. 159–164.
- Djankov, Simeon, Rafael La Porta, Florencio Lopez de Silanes, and Andrei Shleifer**, “Courts,” *The Quarterly Journal of Economics*, 2003, 118, 453–517.
- Donald, Alexia and Florentin Grosset**, “Complementarities in Labor Supply: Joint Commute and Work Decisions,” *Working Paper*, 2024.
- Dunn, Abe, Joshua D. Gottlieb, Adam Hale Shapiro, Daniel J. Sonnenstuhl, and Pietro Tebaldi**, “A Denial a Day Keeps the Doctor Away,” *Quarterly Journal of Economics*, February 2024, 139 (1).
- Earle, John S. and Klara Sabirianova Peter**, “Complementarity and Custom in Wage Contract Violation,” *The Review of Economics and Statistics*, November 2009, 91 (4), 832–849.
- **and Klara Z. Sabirianova**, “How Late to Pay? Understanding Wage Arrears in Russia,” *Journal of Labor Economics*, July 2002, 20 (3), 661–707.
- Falk, Armin**, “Gift Exchange in the Field,” *Econometrica*, 2007, 75 (5), 1501–1511.
- Fernandes-Brough, Jacinta**, “Nigeria Women’s World Cup Players Speak Out Over Unpaid Wages,” September 2023. The Athletic, accessed October 17, 2025.
- Flynn, Suzanne and Mario Pessoa**, “Prevention and Management of Government Expenditure Arrears,” Technical Report TNM/14/03, International Monetary Fund, Fiscal Affairs Department, Washington, DC 2014.

- Freeman, Richard B., Xiaofei Pan, Xiaolan Yang, and Maoliang Ye**, “Team incentives and lower ability workers: A real-effort experiment,” *Journal of Economic Behavior & Organization*, 2025, 233, 106986.
- Galvin, Daniel J.**, “Deterring Wage Theft: Alt-Labor, State Politics, and the Policy Determinants of Minimum Wage Compliance,” *Perspectives on Politics*, 2016, 14 (2), 324–350.
- Gerry, Christopher J., Carmen A. Li, and Byung-Yeon Kim**, “The Gender Wage Gap and Wage Arrears in Russia: Evidence from the RLMS,” *Journal of Population Economics*, 2004, 17 (2), 267–288.
- Guiteras, Raymond P. and B. Kelsey Jack**, “Productivity in piece-rate labor markets: Evidence from rural Malawi,” *Journal of Development Economics*, 2018, 131, 42–61.
- Huck, Steffen, Andrew J. Seltzer, and Brian Wallace**, “Deferred Compensation in Multiperiod Labor Contracts: An Experimental Test of Lazear’s Model,” *American Economic Review*, 2011, 101 (2), 819–843.
- Igwe, Isaac O. C.**, “Rule of Law and Constitutionalism in Nigerian Democracy: A Critical Relativism Discussion in the Context of International Law,” *Athens Journal of Law*, July 2021, 7 (3), 317–334.
- Imbens, Guido W. and Joshua D. Angrist**, “Identification and Estimation of Local Average Treatment Effects,” *Econometrica*, 1994, 62 (2), 467–475.
- Kapoor, Varun**, “Wage theft, reneging and liquidity constraints in informal labor markets,” 2025. Mimeo.
- Kaur, Supreet, Michael Kremer, and Sendhil Mullainathan**, “Self-Control at Work,” *Journal of Political Economy*, 2015, 123 (6), 1227–1277.
- , **Sendhil Mullainathan, Suanna Oh, and Frank Schilbach**, “Do Financial Concerns Make Workers Less Productive?,” *The Quarterly Journal of Economics*, 2 2025, 140 (1), 635–689.
- Krueger, Alan B. and Alexandre Mas**, “Strikes, Scabs, and Tread Separations: Labor Strife and the Production of Defective Bridgestone/Firestone Tires,” *Journal of Political Economy*, 2004, 112 (2), 253–289.
- Laeven, Luc and Christopher Woodruff**, “The Quality of the Legal System, Firm Ownership, and Firm Size,” *The Review of Economics and Statistics*, 2007, 89, 601–614.
- Langener, AM, AW Kramer, W van den Bos, and HM Huizenga**, “A shortened version of Raven’s standard progressive matrices for children and adolescents,” *British Journal of Developmental Psychology*, March 2022, 40 (1), 35–45. Epub 2021 May 27.
- Lazear, Edward P.**, “Why is there mandatory retirement?,” *Journal of Political Economy*, 1979, 87 (6), 1261–1284.
- , “Performance Pay and Productivity,” *American Economic Review*, 2000, 90 (5), 1346–1361.

- Lehmann, Hartmut and Jonathan Wadsworth**, “Wage Arrears and Inequality in the Distribution of Pay: Lessons from Russia,” in Solomon W. Polachek and Olivier Bargain, eds., *Aspects of Worker Well-Being*, Vol. 26 of *Research in Labor Economics*, Leeds: Emerald Group Publishing Limited, 2007, pp. 125–155.
- , – , and **Alessandro Acquisti**, “Grime and Punishment: Job Insecurity and Wage Arrears in the Russian Federation,” *Journal of Comparative Economics*, 1999, 27 (4), 595–617.
- Liu, Licheng, Ye Wang, and Yiqing Xu**, “A Practical Guide to Counterfactual Estimators for Causal Inference with Time-Series Cross-Sectional Data,” *American Journal of Political Science*, 2024, 68 (1), 160–176.
- Mas, Alexandre**, “Pay, Reference Points, and Police Performance,” *Quarterly Journal of Economics*, 2006, 121 (3), 783–821.
- McKenzie, David and Christopher Woodruff**, “Business Practices in Small Firms in Developing Countries,” *Management Science*, 2017, 63, 2967–2981.
- Milkman, Ruth, Ana Luz González, Victor Narro, Annette Bernhardt, Nik Theodore, Douglas Heckathorn, Mirabai Auer, James DeFilippis, Jason Perelshteyn, Diana Polson, and Michael Spiller**, “Wage Theft and Workplace Violations in Los Angeles: The Failure of Employment and Labor Law for Low-Wage Workers,” Technical Report, UCLA Institute for Research on Labor and Employment, Los Angeles, CA 2010.
- National Bureau of Statistics**, “Nigeria Labour Force Survey (NLFS) Report, Q2 2024,” November 2024. Report Date: November 2024.
- Oh, Suanna**, “Does Identity Affect Labor Supply?,” *American Economic Review*, 2023, 113 (8), 2055–2083.
- Punch**, “Oil Workers Protest Over Unpaid Salaries,” 2023.
- PwC**, “MSME Survey 2020 Final,” Survey Report, PricewaterhouseCoopers Nigeria 2020.
- , “MSME Survey 2024: Building resilience – Strategies for MSME success in a changing landscape,” Survey Report, PricewaterhouseCoopers Nigeria 2024.
- Robinson, Erin, Ha T. Nguyen, Scott Isom, Sara A. Quandt, Joseph G. Grzywacz, Haiying Chen, and Thomas A. Arcury**, “Wages, Wage Violations, and Pesticide Safety Experienced by Migrant Farmworkers in North Carolina,” *NEW SOLUTIONS: A Journal of Environmental and Occupational Health Policy*, 2011, 21 (2), 251–268.
- Sánchez de la Sierra, Raúl**, “Whither Formal Contracts?,” *Econometrica*, 2021, 89 (5), 2341–2373.
- , **Kristof Titeca, Haoyang (Stan) Xie, Aimable Amani Lameke, and Albert Malukisa Nkuku**, “The Real State: Inside the Congo’s Traffic Police Agency,” *American Economic Review*, December 2024, 114 (12), 3976–4014.

US Department of State, Bureau of Democracy, Human Rights and Labor, “Country Reports on Human Rights Practices for 2023: Nigeria,” Human Rights Report, United States Department of State 2023.

Vazquez-Bare, Gonzalo, “Identification and estimation of spillover effects in randomized experiments,” *Journal of Econometrics*, 2023, *237* (1), 105237.

Appendices

A Descriptive Survey

A.1 Methodology and Data

A.2 Additional Descriptive Results

We collected data in Lagos between June and August 2025. A team of seven enumerators surveyed 1,279 employees and employers across a range of market settings.⁵⁶ In addition, enumerators visited industrial areas and office complexes to reach workers employed in larger firms. Surveys were administered on tablet computers, and participants were reimbursed with 1,000 NGN (approximately \$0.66). The average survey lasted 21.2 minutes, with a median duration of 21.8 minutes.⁵⁷ Enumerators were assigned to predetermined areas, which they subdivided into sections to minimize overlap and reduce the risk of interviewing the same individual more than once. Each enumerator began at a different location within their section, and the tablet software then instructed them to skip a randomized number of individuals before approaching the next potential participant. Appendix Table A.1 summarizes the characteristics of the respondents. The average participant was 30.8 years old, and 50 percent were female. 32 percent reported actively searching for a job, and the mean monthly income was approximately 199,000 NGN (about \$132). The sample includes both employed (42 percent) and self-employed individuals (56 percent), and spans a wide range of firm sizes. Respondents were also relatively well-educated, mirroring the patterns observed in our recruitee samples.

⁵⁶One of the co-authors also conducted a small subset of surveys.

⁵⁷For 16 surveys, the recorded duration was an implausible 2.18 minutes, with all 16 also showing exactly the same value — an outcome that is statistically highly unlikely. We therefore attribute this to a software error in the *ODK* start and end functions and exclude these cases when calculating the mean and median durations.

Table A.1: Characterisitcs of the Survey Participants

	Mean	SD	10th Percentile	90th Percentile	Observations
Characteristics					
Age	30.81	8.08	22.00	41.00	1,279
% Female	0.50	0.50	0.00	1.00	1,279
% Searching for a Job	0.32	0.46	0.00	1.00	1,279
Income	199,131.32	239,910.47	40,000.00	500,000.00	1,183
Employment Status					
Unemployed	0.02	0.15	0.00	0.00	1,279
Self-Employed	0.42	0.49	0.00	1.00	1,279
Employed	0.56	0.50	0.00	1.00	1,279
At small firm (<10 employees)	0.46	0.50	0.00	1.00	714
At medium firm (10 to 50 employees)	0.30	0.46	0.00	1.00	714
At large firm (>50 employees)	0.24	0.43	0.00	1.00	714
Highest Schooling					
Secondary School	0.53	0.50	0.00	1.00	1,279
Vocational Training	0.00	0.03	0.00	0.00	1,279
National Diploma	0.23	0.42	0.00	1.00	1,279
University	0.23	0.42	0.00	1.00	1,279

NOTE: This Table shows the characteristics of the survey participants.

A.2.1 Additional Descriptive Results

We investigate individuals' concerns for working at different types of firms. We distinguish between large and small firms. The exact wording used in the survey for large firms is "Imagine you work for a large and established company with a written work agreement. Would you be concerned about." For small firms, we wanted to distinguish between firms owned by a friend or family member, and by someone not necessarily from the respondent's social network.

Table A.2: Firm Concerns

	Mean Response			Difference	
	Large Firm	Small Firm (non-family owned)	Small Firm (family owned)	Small firms family - non-family owned	Large Firm-small firm
Concern: being fired and without job	0.20	0.18	0.15	0.03 (0.021)	0.02 (0.021)
Concern: paid as agreed and on time	0.31	0.50	0.45	0.04 (0.028)	-0.19*** (0.024)
Concern: employer sticking to work agreement	0.21	0.27	0.33	-0.05** (0.026)	-0.07*** (0.020)
Concern: poorly treated b/c of ethnicity	0.03	0.05	0.02	0.03** (0.011)	-0.01 (0.009)
Concern: poorly treated b/c of religion	0.03	0.05	0.02	0.03** (0.010)	-0.02** (0.007)
No Concern	0.40	0.26	0.31	-0.06** (0.025)	0.15*** (0.022)
Observations	1,279	654	625	1,279	654

NOTE: This Table shows individuals' concerns for working at different types of firms. We distinguish between large and small firms. The exact wording used in the survey is "Imagine you work for a large and established company with a written work agreement."

Table A.3: Coping with Salary Non-Payment and Its Extent within the Firm

	Mean	SD	Observations
Panel A: Coping with Delayed Salaries			
Borrowed from family/friends	0.31	0.46	383
Borrowed from moneylender	0.04	0.20	383
Had Savings	0.45	0.50	383
Had other income	0.33	0.47	383
Panel B: Wage Withholding Among Employees			
Participant was the only unpaid employee	0.23	0.42	383
Some employees were unpaid	0.46	0.50	383
All employees were unpaid	0.31	0.46	383

NOTE: TBD

Table A.4: Coping with Salary Non-Payment and Its Extent within the Firm

	Mean	SD	Observations
Panel A: Social Norm of Wage Payments			
Employers should pay their workers	0.70	0.46	1,279
Employers may delay payments if in trouble	0.22	0.42	1,279
Workers should temporarily forgo wages if employer is in trouble	0.07	0.26	1,279
Panel B: Wage Withholding in Social Network			
Knows someone who experienced salary difficulties	0.49	0.50	1,279

NOTE: TBD

Table A.5: Additional Beliefs about Salary Payments

	Mean	SD	Observations
Perceived Risk of Salary Difficulty			
At new firm	0.44	0.28	1,271
At reputable firm (Dangote)	0.20	0.21	1,279

NOTE: TBD

Table A.6: Preferences for Self-Employment

	Mean	SD	10th Percentile	90th Percentile	Observations
Preference for Self-Employment	0.82	0.38	0.00	1.00	1,279
Reasons for Self-Employment Preference					
Pref. for self-empl: can grow business and make more money	0.47	0.50	0.00	1.00	1,049
Pref. for self-empl: want to be one's own boss	0.39	0.49	0.00	1.00	1,049
Pref. for self-empl: more stable income than employment	0.33	0.47	0.00	1.00	1,049
Pref. for self-empl: Concerned about employer sticking to work agreement	0.07	0.26	0.00	0.00	1,049
Pref. for self-empl: Concerned to be fired and without job	0.09	0.28	0.00	0.00	1,049
Pref. for self-empl: employment locks in, good opportunity can be missed	0.09	0.28	0.00	0.00	1,049
Pref. for self-empl: self-employment is status symbol	0.07	0.25	0.00	0.00	1,049

NOTE: TBD

B Applicants

B.1 In-Person Recruitment

In-person recruitment was carried out by a team of eight enumerators. For the recruitment task, they were formally affiliated with the recruitment agency *Unlocking Creativity* and wore clothing and ID cards displaying the agency’s branding.

Enumerators recruited participants in pre-selected markets and public areas. Each was equipped with a tablet computer running *ODK Collect* to conduct the survey. The survey form implemented a pre-specified randomization protocol. Before approaching an individual, enumerators followed a skip pattern that required them to pass over a randomly determined number of suitable persons (between one and ten). The software also randomly specified whether to approach a man or a woman (each with 50 percent probability) and, independently, whether to approach someone who appeared to be self-employed with a small business (60 percent probability) or someone who was not self-employed and appeared to be passing by (40 percent probability). Appendix Figure B.3 illustrates the number of interactions of the enumerators per recruitment day.

Table B.7: Balance of Baseline Sample

	Treatment Arm 1		Treatment Arm 2		
	Mean	SD	Mean	SD	Difference
Characteristics					
Age	29.48	10.21	31.31	9.76	-1.82 (1.868)
% Female	0.55	0.51	0.54	0.50	0.01 (0.093)
% Searching for a Job	1.00	0.00	0.99	0.08	0.01 (0.005)
Ravens Score (0–15)	8.82	3.91	7.93	4.27	0.88 (0.725)
Employment Status					
Unemployed	0.36	0.49	0.46	0.50	-0.09 (0.090)
Employed	0.06	0.24	0.10	0.30	-0.04 (0.046)
Self-Employed	0.58	0.50	0.44	0.50	0.13 (0.092)
Highest Schooling					
Secondary School	0.45	0.51	0.33	0.47	0.12 (0.092)
Vocational Training	0.03	0.17	0.01	0.12	0.02 (0.031)
National Diploma	0.09	0.29	0.16	0.37	-0.07 (0.055)
University	0.42	0.50	0.49	0.50	-0.07 (0.092)
Observations	33	33	288	288	321

NOTE: TBD

Table B.8: Recruitment Sample

	Mean	SD	10th Percentile	90th Percentile	Observations
Panel A: In-Person Recruitment Sample (Incentivized)					
Characteristics					
Age	27.16	7.57	18.00	38.00	821
% Female	0.42	0.49	0.00	1.00	821
Raven's Matrices	7.40	3.80	2.00	13.00	821
% Searching for a Job	0.62	0.49	0.00	1.00	821
% Attending Orientation	0.44	0.50	0.00	1.00	821
Employment Status					
Unemployed	0.26	0.44	0.00	1.00	821
Employed	0.15	0.36	0.00	1.00	821
Self-Employed	0.59	0.49	0.00	1.00	821
Highest Schooling					
Secondary School	0.63	0.48	0.00	1.00	821
Vocational Training	0.01	0.12	0.00	0.00	821
National Diploma	0.13	0.33	0.00	1.00	821
University	0.20	0.40	0.00	1.00	821
Panel B: In-Person Recruitment Sample (Unincentivized)					
Characteristics					
Age	25.65	7.12	18.00	35.00	258
% Female	0.40	0.49	0.00	1.00	258
Raven's Matrices	8.09	4.27	2.00	13.00	258
% Searching for a Job	0.89	0.32	0.00	1.00	258
% Attending Orientation	0.60	0.49	0.00	1.00	258
Employment Status					
Unemployed	0.47	0.50	0.00	1.00	258
Employed	0.11	0.32	0.00	1.00	258
Self-Employed	0.41	0.49	0.00	1.00	258
Highest Schooling					
Secondary School	0.60	0.49	0.00	1.00	258
Vocational Training	0.02	0.14	0.00	0.00	258
National Diploma	0.15	0.36	0.00	1.00	258
University	0.22	0.42	0.00	1.00	258
Panel C: Job Advertisement Sample					
Characteristics					
Age	29.53	9.38	18.00	42.00	638
% Female	0.53	0.50	0.00	1.00	638
Raven's Matrices	8.16	4.15	2.00	13.00	638
% Searching for a Job	0.99	0.09	1.00	1.00	638
% Attending Orientation	0.95	0.22	1.00	1.00	638
Employment Status					
Unemployed	0.51	0.50	0.00	1.00	638
Employed	0.10	0.30	0.00	0.00	638
Self-Employed	0.39	0.49	0.00	1.00	638
Highest Schooling					
Secondary School	0.39	0.49	0.00	1.00	638
Vocational Training	0.01	0.09	0.00	0.00	638
National Diploma	0.16	0.37	0.00	1.00	638
University	0.44	0.50	0.00	1.00	638

NOTE: This table shows average characteristics and statistics across our three different samples. Sample I: in-person recruited and requiring an incentive to participate in the job information session. Sample II: in-person recruited and willing to participate in the job information session without an incentive. Sample III: job advertisement sample, i.e. individuals who responded to our job advertisement.

Table B.9: Employee Sample

	Mean	SD	10th Percentile	90th Percentile	Observations
Panel A: Job Offer Acceptance					
Job offer accepted (Control Offer 1)	1.00	0.00	1.00	1.00	33
Job offer accepted (Uncertainty Offer 2)	0.95	0.22	1.00	1.00	281
Panel B: Employee Demographics					
Characteristics					
Age	30.72	9.73	20.00	44.50	300
% Female	0.54	0.50	0.00	1.00	300
% Searching for a Job	0.99	0.08	1.00	1.00	300
Employment Status					
Unemployed	0.46	0.50	0.00	1.00	300
Employed	0.08	0.27	0.00	0.00	300
Self-Employed	0.46	0.50	0.00	1.00	300
Highest Schooling					
Secondary School	0.35	0.48	0.00	1.00	300
Vocational Training	0.02	0.13	0.00	0.00	300
National Diploma	0.15	0.36	0.00	1.00	300
University	0.48	0.50	0.00	1.00	300
Panel C: Employee Work Patterns					
Total Treatment Assignments	1.03	0.98	0.00	3.00	300
Absent (0,1)	0.21	0.20	0.02	0.50	300
Hours at work	7.16	0.32	6.89	7.38	300
Time spent working	6.12	0.58	5.40	6.71	300
Images labeled per day	255.73	156.71	85.68	487.99	300
Images labeled in total	15,332.38	9,409.85	5,130.00	29,275.00	300
Time spent per image	89.96	50.73	41.88	156.16	300
Correct labels per image (0,10)	8.37	0.39	7.97	8.66	300
All labels correct (0,1)	0.10	0.04	0.04	0.14	300

NOTE: This table shows the acceptance rate of job offers and the characteristics of the employee sample.

Figure B.1: Job Advertisement

VACANCY! VACANCY!! VACANCY!!!

**DATA CLASSIFICATION WORKERS ARE
NEEDED WITH NO SPECIAL
COMPUTER SKILLS ARE REQUIRED.**

SALARY: 50K -85K MONTHLY

**CONTACT THIS NUMBER
Via Call**

09122018004

OR WhatsApp

08105770268


NOTE: This figure shows the job advertisement that we pasted across Lagos. Lego images.

Figure B.2: Job Advertisement




NOTE: This figure shows the job advertisement that we pasted across Lagos. Lego images.

Figure B.3: In-Person Recruitment Timing



NOTE: This figure shows the timeline of our in-person recruitment. The teal darker teal bars show the number of interactions of our enumerators that were successful and lead to a job-information session. The lighter orange bars show the number of interactions that did not lead to a job-information session. On average, 44 percent of interactions were successful (mean = 0.44, SD = 0.50).

Figure B.4: Interview Timing



NOTE: This figure shows the timeline of job interviews with jobseekers who responded to our job advertisements. Bars indicate the number of interviews conducted per day, restricted to days on which interviews took place

C Image Labelling Task

Lego pieces are plastic building blocks made for children to play with that come in many shapes, sizes, and colors. They are designed to fit together securely, allowing the creation of all kinds of complex structures.

C.1 Creation of Images

All images used in the labeling task were generated digitally using the Lego bricks online simulator *mecabricks.com*. This simulator provides access to the complete set of Lego brick types and allows users to arrange them in a virtual 3D space. The resulting brick geometries can be exported and processed, for example, via a Python script to produce images.

We used the simulator to create 1,000 unique Lego brick geometries: 100 containing a single brick and 900 containing multiple bricks. Each geometry was imported into the image renderer *Autodesk Maya 2024*, where we generated 40 images per geometry with a plain white background, varying the camera position for each shot. This procedure generated 40,000 unique images. From this set, we randomly selected 4,000 images and replaced the plain white background with a visually noisy, captcha-style background. This resulted in a final dataset of 44,000 unique images: 40,000 regular images and 4,000 ‘high-stakes’ images for our labelling task.

Because the images were fully computer-generated, we maintained complete control over attributes such as color, spatial arrangement, and viewing angle. By generating all images digitally ourselves, we maintained complete control over every attribute—including color, spatial arrangement, and viewing angle. This allowed us to define the true classification for each image and evaluate employees’ performance against it.

C.2 Labelling Categories

Employees were required to label each image according to six categories of heterogeneous complexity: (a) the color of the Lego pieces (all pieces in a given image were the same color); (b) whether any Lego pieces in the image were stacked together; (c) the total number of individual Lego pieces in the image; (d) the types of Lego pieces visible in the image (bricks, plates, bows, circles, and angles); (e) the type of Lego piece appearing most frequently in the image; and (f) whether a 2×2 Lego brick was visible in the image.

(a) The Color of the Lego Pieces. All Lego pieces in a given image were the same color. Employees were required to identify this color and select the correct option from a dropdown list of nine predefined colors: red, orange, yellow, green, teal, blue, purple, brown, and black. A response was classified as correct if the chosen color matched the color assigned to the bricks during the generation of the digital image.

(b) Whether any Lego Pieces in the Image Were Stacked Together. In most images, Lego pieces were arranged without touching each other. However, in some geometries — and consequently in some images — some or all pieces were arranged so that they appeared to be stacked together. Employees were required to determine whether any pieces in a given image were stacked and to select the correct “yes” or “no” response from a dropdown list to the question: “Are the pieces in the image stacked together?”

(c) Total Number of Individual Lego Pieces in the Image. Images contained between one and seven individual Lego pieces. Employees were required to count the number of pieces and select the correct value from a dropdown list of options ranging from one to seven.

(d) Types of Lego Pieces Visible in the Image. All pieces included in the images belonged to one of five categories: bricks, plates, bows, circles, or angles. Bricks were

defined as any cubic piece. Plates were defined as any flat rectangular piece. Bows were defined as curved pieces that are not circular. Circles were defined as any circular piece. Angles were defined as any piece without curves and with at least one angle different from 90° .

An image could contain pieces from one or multiple categories. Employees were required to tick a checkbox for each type visible in the image; leaving a checkbox unticked indicated that the type was not visible. This yielded five separate responses — one for each category. For each category, we compared the employee’s response with the correct classification for that image. A response was classified as correct if the checkbox selection matched the true appearance of the respective category in the given image.

(e) Type of Lego Piece Appearing Most Frequently in the Image. Employees were required to identify the type of piece that appeared most often in the image. For example, if the image contained three bricks, two plates, and one circle, the correct response would be “brick.” Employees selected the answer from the list of five possible types.

If two or more types appeared the same number of times — for instance, two plates and two circles — employees were required to select the option “no single type appears most often” instead of choosing one of the five types. A response was classified as correct if the selected option from the dropdown list matched the true most-frequent type or correctly indicated that no single type appeared most often.

(f) Whether a 2×2 Lego Brick Was Visible in the Image. Employees were asked to indicate whether a specific Lego piece—a 2×2 brick—was visible in the image. This is a common Lego piece, illustrated in Appendix Figure C.1, and is relatively easy to identify. Employees selected either “yes” or “no” from a dropdown list in response to the question: “Is there a 2×2 brick in the image?”

C.3 Set up of the task.

All employees were provided with identical laptops for the task, as well as individual *Labelbox* accounts with unique login credentials that we set up in advance. Once logged in, they could begin labeling immediately. Each image was shown only once, and employees had no opportunity to return to a previously submitted image to change their answers. No feedback on performance — either absolute or relative to other employees — was provided within *Labelbox* or in any other form. Appendix Figure C.3 shows the *Labelbox* labeling interface with an example image, presented exactly as employees saw it during the task.

On the first day, employees received extensive training on the image classification task and the use of *Labelbox*. Each employee had a PDF copy of the training materials on their laptop and was encouraged to consult it whenever they had questions about a labeling task. In addition, each work location had physical examples of Lego pieces available, so employees could familiarize themselves with the different types of pieces.

C.4 Representativeness of the Lego Task

The Lego task exemplifies *data labeling*, a core activity in modern AI development. Data labeling is the process of adding descriptive information to raw data —for example, classifying objects in images or identifying the sentiment of a text. These labels provide the examples that machine-learning systems need in order to learn and make predictions about new observations. Because algorithms cannot reliably generate such labels on their own, data labeling is typically carried out by human workers around the world.

For instance, crowdsourced workers across the globe (e.g., on platforms like Appen or Scale AI) label images, text, and video for major tech firms, often in low-income countries (Chen, 2023; *Humans Are Essential for Training AI*, 2024; *The Hidden Workforce Behind Artificial Intelligence*, 2023). Kassi et al. (2021) estimate that over 163 million freelancer profiles are registered on online labor platforms, with approximately 19 million having per-

formed at least one job, and 5 million completing at least ten tasks or earning over USD 1,000. While not all of this work involves data labeling, Appen alone reports over one million data-labeling contributors globally (Chen, 2023). The growing institutionalization of this sector is also evident in the creation of organizations such as the *Data Labeling Association of Kenya*, which represents local workers engaged in annotation tasks and advocates for fair labor standards.⁵⁸

Data labeling is among the fastest-growing forms of digital work globally, with the market for labeling solutions projected to grow at more than 20% annually through 2030 (G2 Learning Hub, 2023). The broader AI data-labeling market is forecast to reach USD 134.7 billion by 2034, up from USD 19.7 billion in 2024 (Market.us Insights, 2024).

This scale illustrates that data labeling is a ubiquitous form of digital labor. Through the Lego-image labeling task, we reproduce the core elements of large-scale labeling workflows — repetition, cognitive engagement, and precision in categorization — within a controlled and measurable setting. These features also characterize many other forms of work.


C.5 Image-Level Summary Statistics

Table C.1 summarizes the resulting dataset of 4,595,228 labelled images. Panel (a) presents data at the image-employee level, with each observation representing a single employee’s labeling of an image. Employees spent an average of 68.4 seconds labeling each image and correctly assigned an average of 8.4 labels per image. Overall, this translates into an 11 percent probability that an image is labeled entirely correctly by any given employee, indicating considerable variation in labeling accuracy and speed.

Panel (b) aggregates data to the image level, with each of the 44,000 unique images representing a single observation, averaged across multiple employee labels. At this aggregated level, the average accuracy per image remains 11 percent, with each image labeled by an average of 104 different employees.


⁵⁸See: <https://datalabelers.org>

Figure C.1: 2×2 Brick




NOTE: This figure shows a 2×2 brick.

Figure C.2: Image with Captcha Style Background



NOTE: This figure shows an example image from the labeling task with a visually noisy, captcha-style background. Such images were designated as ‘high-stakes’, and employees were instructed to devote extra effort and care when labeling them.

Figure C.3: Examples Image of Lego Bricks



NOTE: This figure shows a screenshot of an example Lego bricks image with the entry boxes for the classification of the image as employees have seen during the experiment. This screenshot is taken from the data classification platform *Labelbox*, which was used during the field experiment to classify the Lego images.

Table C.1: Lego Bricks Data

	Mean	SD	10th Percentile	90th Percentile	Observations
Panel A: Image Employee Level Dataset					
Correct Labels (0–10)	8.43	1.01	7.00	10.00	4,595,228
All labels correct (0,1)	0.11	0.31	0.00	1.00	4,595,228
Time spent per image	68.44	115.78	16.00	137.00	4,595,228
Panel B: Image Level Dataset					
Correct Labels (0–10)	8.53	0.75	7.67	9.67	44,000
All labels correct (0,1)	0.11	0.27	0.00	0.74	44,000
Time spent per image	47.93	24.13	18.25	80.81	44,000
Employees	104.44	115.33	4.00	296.00	44,000

NOTE: This table summarizes the dataset of Lego images and corresponding labelling performance from the first employment round (the second employment round is ongoing at the moment).

D Validity Checks

$$Y_{itq} = \alpha + \delta_q + \kappa_i + \sigma_t + \sum_{k=T-1}^{-4} \beta_k \times \text{treat}_{ik} + \varepsilon_{itq} \quad (12)$$

E Robustness

E.1 Alternative Definition of Productivity Measure

In this subsection, we examine labeling performance and treatment effects using a pre-registered subset of the labeling questions. The index is constructed from questions (c) and (d) described in Appendix C.2. Accordingly, the continuous index ranges from zero to six, while the binary index equals one if all six questions are answered correctly and zero otherwise. Table E.2 replicates Table 3 — our main analysis table — using the alternative productivity index. Table E.3 replicates Table E.15 — the table showing effects for the most difficult images — using the alternative productivity measure. Qualitative results do not change when using the alternative native productivity measure.

Table E.2: Treatment Effects on Worker Effort for Alternative Productivity Measure

	All Images		High-Stakes Images	
	Continuous Index	Binary Index	Continuous Index	Binary Index
ATE	0.0138 (0.0043)*** [0.0046]***	0.0013 (0.0006)** [0.0007]**	0.0147 (0.0057)** [0.0071]**	0.0004 (0.0012) [0.0014]
Observations (First Stage)	2,166,877	2,166,877	196,994	196,994
Observations (Imputed)	675,184	675,184	61,209	61,209
Individuals	297	297	297	297
SE Cluster	Individual	Individual	Individual	Individual
Mean of Dep. Var.	4.69	.14	4.68	.13
Q-Value	0.005	0.025	0.011	0.235

NOTE: This table shows the effect of being owed salary payments on employees' effort measures. The treatment effects are estimated at the worker-image level using the Borusyak et al. (2024) imputation estimator. The daily treatment effects are combined into a single parameter spanning a pay cycle or a week using the delta method. Standard errors are initially clustered at the individual level and shown in parentheses. Since we measure productivity in different ways (with a continuous and binary index) we also report sharpened q -values in brackets to account for multiple hypotheses testing.

Table E.3: Effect of Salary Delay on Workers' Effort (Hard Images and Alternative Indices)

	Most Difficult 10 Percent		Most Difficult 25 Percent		Most Difficult 50 Percent	
	Continuous Index	Binary Index	Continuous Index	Binary Index	Continuous Index	Binary Index
ATE	0.0121 (0.0065)* [0.0081]*	0.0000 (0.0004) [0.0005]	0.0160 (0.0045)*** [0.0057]***	0.0001 (0.0002) [0.0002]	0.0160 (0.0050)*** [0.0077]***	-0.0001 (0.0002) [0.0002]
Observations (First Stage)	199,276	199,276	496,231	496,231	975,265	975,265
Observations (Imputed)	42,053	42,053	114,484	114,484	230,584	230,584
Individuals	297	297	297	297	297	297
SE Cluster	Individual	Individual	Individual	Individual	Individual	Individual
Mean of Dep. Var.	3.652	.001	3.831	.001	4.168	.003
Q-Value	0.100	0.933	0.002	0.237	0.006	0.561

NOTE: tbd

E.2 Spillover Concerns

A key threat to the experimental design is the possibility of spillover effects because employees with different treatment status work in the same location. In this setting, morale or other behavioral responses may spread from treated to untreated employees, potentially biasing the estimated treatment effects. We address this threat in two ways. First, we use outcomes from a location that was deliberately assigned to contain only untreated employees to impute counterfactual outcomes for treated employees. Second, within mixed-treatment locations, we explicitly designed fixed seating arrangements to create reference groups, allowing us to account for potential spillover effects within each group.

Mitigating Spillovers with an Untreated-Only Location. The first way we address potential spillover effects is by relying on untreated employees from a work location with only untreated workers. We replicate Table 3 and Figure 4, restricting the sample of untreated employees to this group. The key advantage of this approach is that these employees are certainly unaffected by treated peers, since they work in a separate location. Two disadvantages remain. First, we still rely on untreated pre-periods of employees who share a location with other treated workers, so some concerns about spillovers may remain. Second, restricting the sample of untreated employees to the ones working in the untreated-only location reduces sample size and statistical power, making the estimates noisier. Nonetheless, the results are qualitatively similar to those reported in Table 3 and Figure 4. Table E.4 presents estimates using the standard continuous and binary indices, while Table E.5 reports results with the alternative indices. Columns1-2 reproduce the estimates based on all images for reference, and Columns3-4 restrict to untreated observations from the untreated-only location to impute counterfactual outcomes for the treated sample.

Table E.4: Treatment Effects: Imputations Using Untreated-Only Location

	All Images		Untreated-Only Location	
	Continuous Index	Binary Index	Continuous Index	Binary Index
ATE	0.0449 (0.0164)*** [0.0176]***	0.0062 (0.0027)** [0.0029]**	0.0330 (0.0138)** [0.0185]**	0.0030 (0.0019) [0.0029]
Observations (First Stage)	2,166,877	2,166,877	1,368,152	1,368,152
Observations (Imputed)	675,184	675,184	675,184	675,184
Individuals	297	297	217	217
SE Cluster	Individual	Individual	Individual	Individual
Mean of Dep. Var.	8.43	.11	8.44	.11
Q-Value	0.010	0.016	0.026	0.047

NOTE: This table reports the effect of being owed salary payments on employees' effort measures. Columns 1-2 replicate the results from Table 3, while Columns 3-4 use untreated employees from the untreated-only location to impute counterfactual outcomes for treated employees. Treatment effects are estimated at the worker-image level using the Borusyak et al. (2024) imputation estimator. Daily treatment effects are aggregated into a single parameter over a pay cycle or week using the delta method. Standard errors, clustered at the individual level, are shown in parentheses. Because productivity is measured in two ways (continuous and binary indices), we report sharpened q -values in brackets to account for multiple hypothesis testing.

Table E.5: Treatment Effects: Imputations Using Untreated-Only Location (Alternative Productivity Measure)

	All Images		Untreated-Only Location	
	Continuous Index	Binary Index	Continuous Index	Binary Index
ATE	0.0138 (0.0043)*** [0.0046]***	0.0013 (0.0006)** [0.0007]**	0.0169 (0.0043)*** [0.0057]***	0.0016 (0.0007)** [0.0009]**
Observations (First Stage)	2,166,877	2,166,877	1,368,152	1,368,152
Observations (Imputed)	675,184	675,184	675,184	675,184
Individuals	297	297	217	217
SE Cluster	Individual	Individual	Individual	Individual
Mean of Dep. Var.	4.69	.14	4.69	.14
Q-Value	0.005	0.025	0.001	0.035

NOTE: This table reports the effect of being owed salary payments on employees' effort measures using alternative productivity measures. Columns 1-2 replicate the results from Table E.2, while Columns 3-4 use untreated employees from the untreated-only location to impute counterfactual outcomes for treated employees. Treatment effects are estimated at the worker-image level using the Borusyak et al. (2024) imputation estimator. Daily treatment effects are aggregated into a single parameter over a pay cycle or week using the delta method. Standard errors, clustered at the individual level, are shown in parentheses. Because productivity is measured in two ways (continuous and binary indices), we report sharpened q -values in brackets to account for multiple hypothesis testing."

Mitigating Spillovers with Designed Reference Groups. Our second approach to addressing potential spillover effects is to create a clearly defined reference group for each employee, which allows us to control for and estimate spillovers within a work location. To construct these reference groups credibly within an open-space work location, we relied on three design features: fixed seating arrangements, restrictions on communication across groups, and staggered work schedules. First, employees were seated in groups of four around a single table, with fixed seats that remained unchanged throughout the employment period, ensuring that each individual interacted with the same peers for the duration of the job. Second, we restricted communication across tables by prohibiting conversation during work, generally allowing only one person at a time to use the washroom, and providing bottled water at each table to eliminate the need for a shared water fountain. Third, we implemented staggered start, end, and break times across tables to prevent crowding and mingling when employees arrived or left work and during break periods. We provide evidence that these measures successfully created clearly defined reference groups in Figure E.4.

Panel (a) shows that there was virtually no communication across tables, based on our daily supervisory staff survey (see Appendix G for survey details). Panel (b) demonstrates that employees were well acquainted with their table peers but not with workers at other tables, as measured in our Peer Recognition survey (see Appendix G for survey details). Together, these patterns suggest that the peer groups functioned as intended and that spillovers across groups were minimal.

We then use the reference groups to explicitly control for potential spillover effects. We continue to rely on the imputation estimator of Borusyak et al. (2024), which requires one key assumption about the nature of spillovers. Because the estimator uses untreated observations to estimate fixed effects, we must assume that spillover effects from treated employees affect untreated and treated co-workers in the same way. In this setting, that assumption is reasonable: when working alongside a peer who experiences unpaid salaries, it is plausible

that the effect on others does not depend on whether those others are themselves treated.

Following the notation introduced in Section 6.1, we adapt the approach of Vazquez-Bare (2023) to set up the estimation of spillover effects. Specifically, we estimate the following modified version of equation (5):

$$y_{itq} = \alpha + \delta_q + \kappa_i + \sigma_t + \lambda \mathbb{1}(S_{igt} > 0) + \varepsilon_{itq} \quad (13)$$

via OLS on $(i, t) \in \mathcal{S}_0$ (i.e., untreated observations only). The parameters δ_q , κ_i , and σ_t denote vectors of fixed effects as before. We additionally include S_{igt} , an indicator equal to one if individual i in reference group g is exposed to treated individuals at time t . This can include any number of treated co-workers (see Vazquez-Bare (2023) for further discussion). As before, we then use the estimated parameters from equation 13 to construct counterfactual outcomes for treated individuals. The only difference is that our model for y_{itq} now explicitly accounts for spillover effects. This adjustment allows us to distinguish the direct impact of salary delays on treated workers from indirect effects transmitted through their peers, under the assumption that spillover effects are homogeneous across untreated and treated employees.

Table E.6: Treatment Effects: Clustering at the Table-Level

	All Images		High-Stakes Images	
	Continuous Index	Binary Index	Continuous Index	Binary Index
ATE	0.045 (0.017)** [0.020]**	0.006 (0.003)** [0.003]**	0.048 (0.016)*** [0.019]***	0.008 (0.003)*** [0.004]***
Observations (First Stage)	2,166,877	2,166,877	196,994	196,994
Observations (Imputed)	675,184	675,184	61,209	61,209
Individuals	297	297	297	297
SE Cluster	Table	Table	Table	Table
Mean of Dep. Var.	8.43	.11	8.41	.11
Q-Value	0.016	0.022	0.011	0.011
Spillover Control	No	No	No	No

NOTE: This table reports the effect of being owed salary payments on employees' effort measures from equation (6), aggregated into a single ATE. This table reports an alternative clustering level for the standard errors. Standard errors based on the conservative variance estimator of Borusyak et al. (2024), clustered at the table level and aggregated using the delta method, are shown in parentheses. Bootstrap standard errors with 500 replications, following Liu et al. (2024), are reported in brackets. The standard errors are now clustered at the table-level, i.e. the level of the created reference group. Hence, this table replicates Table 3, clustering standard errors at the table-level.

Table E.7: Treatment Effects: Clustering at the Table-Level and Accounting for Spillover Effects

	All Images		High-Stakes Images	
	Continuous Index	Binary Index	Continuous Index	Binary Index
ATE	0.044 (0.017)** [0.019]**	0.006 (0.003)** [0.003]**	0.047 (0.016)*** [0.021]***	0.007 (0.003)*** [0.004]***
Observations (First Stage)	2,166,877	2,166,877	196,994	196,994
Observations (Imputed)	675,184	675,184	61,209	61,209
Individuals	297	297	297	297
SE Cluster	Table	Table	Table	Table
Mean of Dep. Var.	8.43	.11	8.41	.11
Q-Value	0.017	0.022	0.011	0.011
Spillover Control	Yes	Yes	Yes	Yes

NOTE: This table reports the effect of being owed salary payments on employees' effort measures from equation (13), aggregated into a single ATE. This table presents estimates that account for spillover effects among employees. Standard errors based on the conservative variance estimator of Borusyak et al. (2024), clustered at the table level and aggregated using the delta method, are shown in parentheses. Bootstrap standard errors with 500 replications, following Liu et al. (2024), are reported in brackets. The standard errors are clustered at the table-level, i.e. the level of the created reference group. Hence, this table replicates Table 3, clustering standard errors at the table-level and accounting for spillover effects.

Table E.8: Treatment Effects: Clustering at the Table-Level (Alternative Productivity Measure)

	All Images		High-Stakes Images	
	Continuous Index	Binary Index	Continuous Index	Binary Index
ATE	0.0138 (0.004)*** [0.005]***	0.0013 (0.001)** [0.001]**	0.0147 (0.006)*** [0.007]***	0.0004 (0.001) [0.001]
Observations (First Stage)	2,166,877	2,166,877	196,994	196,994
Observations (Imputed)	675,184	675,184	61,209	61,209
Individuals	297	297	297	297
SE Cluster	Table	Table	Table	Table
Mean of Dep. Var.	4.69	.14	4.68	.13
Q-Value	0.003	0.022	0.011	0.222
Spillover Control	No	No	No	No


NOTE: This table reports the effect of being owed salary payments on employees' effort measures from equation (6), aggregated into a single ATE. This table reports an alternative clustering level for the standard errors. Standard errors based on the conservative variance estimator of Borusyak et al. (2024), clustered at the table level and aggregated using the delta method, are shown in parentheses. Bootstrap standard errors with 500 replications, following Liu et al. (2024), are reported in brackets. The standard errors are now clustered at the table-level, i.e. the level of the created reference group. Hence, this table replicates Table E.2, clustering standard errors at the table-level.

Table E.9: Treatment Effects: Clustering at the Table-Level and Accounting for Spillover Effects (Alternative Productivity Measure)

	All Images		High-Stakes Images	
	Continuous Index	Binary Index	Continuous Index	Binary Index
ATE	0.0139 (0.004)*** [0.005]***	0.0013 (0.001)** [0.001]**	0.0148 (0.006)*** [0.007]***	0.0005 (0.001) [0.001]
Observations (First Stage)	2,166,877	2,166,877	196,994	196,994
Observations (Imputed)	675,184	675,184	61,209	61,209
Individuals	297	297	297	297
SE Cluster	Table	Table	Table	Table
Mean of Dep. Var.	4.69	.14	4.68	.13
Q-Value	0.003	0.022	0.011	0.197
Spillover Control	Yes	Yes	Yes	Yes

NOTE: This table reports the effect of being owed salary payments on employees' effort measures from equation (13), aggregated into a single ATE. This table presents estimates that account for spillover effects among employees. Standard errors based on the conservative variance estimator of Borusyak et al. (2024), clustered at the table level and aggregated using the delta method, are shown in parentheses. Bootstrap standard errors with 500 replications, following Liu et al. (2024), are reported in brackets. The standard errors are clustered at the table-level, i.e. the level of the created reference group. Hence, this table replicates Table E.2, clustering standard errors at the table-level and accounting for spillover effects.

Figure E.4: Reference Group Validation



NOTE: This figure provides suggestive evidence validating our construction of the reference groups. Panel (a) displays Likert scale responses from the daily supervisory staff survey on communication across tables during a given workday, ranging from “None” (1) to “Very Much” (6). The panel also plots a kernel density estimate using an Epanechnikov kernel with a bandwidth of five days, based on the numeric values of the Likert scale. Panel (b) reports employees’ acquaintance with co-workers. The first two bars show employees’ knowledge about the treatment status of their table peers (whether they were currently owed salary or had been owed in the past), with a large share indicating familiarity. In contrast, just a very small share reported knowing the treatment status of a randomly selected co-worker seated at another table (bars three and four), and the vast majority even reported not knowing this co-worker at all (bar five). We also display 95 percent confidence intervals around the sample means.

E.3 Alternative Weighting

In this subsection, we examine labeling performance and treatment effects using an alternative weighting scheme when aggregating image-employee-level treatment effects. Specifically, we weight each image inversely proportional to the number of images labeled by that employee on a given workday. This approach gives equal weight to each employee, rather than to each individual labeled image.

Table E.10: Treatment Effects on Worker Effort for Alternative Weighting

	All Images		High-Stakes Images	
	Continuous Index	Binary Index	Continuous Index	Binary Index
ATE	0.0389 (0.0198)** [0.0176]**	0.0055 (0.0026)** [0.0023]**	0.0347 (0.0210)* [0.0189]*	0.0071 (0.0032)** [0.0031]**
Observations (First Stage)	2,166,877	2,166,877	196,994	196,994
Observations (Imputed)	675,184	675,184	61,209	61,209
Individuals	297	297	297	297
SE Cluster	Individual	Individual	Individual	Individual
Mean of Dep. Var.	8.43	.11	8.41	.11
Q-Value	0.071	0.071	0.119	0.119

NOTE: This table reports the effect of being owed salary payments on employees' effort measures from equation (6), aggregated into a single ATE using alternative weights that are inversely proportional to the number of images labelled by a given employee on a given day. Standard errors based on the conservative variance estimator of Borusyak et al. (2024), clustered at the individual level and aggregated using the delta method, are shown in parentheses. Bootstrap standard errors with 500 replications, following Liu et al. (2024), are reported in brackets. 'First-stage observations' refer to the number of observations used to estimate equation (5), while 'imputed observations' indicate the number of counterfactual outcomes generated to estimate treatment effects. Q-values accounting for multiple hypothesis testing across our four productivity measures are reported, based on the p-values in parentheses. This table replicates Table 3 with different weights.

Table E.11: Effect of Salary Delay on Workers' Effort (Alternative Indices and alternative weighting)

	All Images		High-Stakes Images	
	Continuous Index	Binary Index	Continuous Index	Binary Index
ATE	0.0141 (0.0062)** [0.0052]**	0.0012 (0.0008) [0.0007]	0.0141 (0.0087) [0.0083]	0.0002 (0.0018) [0.0016]
Observations (First Stage)	2,166,877	2,166,877	196,994	196,994
Observations (Imputed)	675,184	675,184	61,209	61,209
Individuals	297	297	297	297
SE Cluster	Individual	Individual	Individual	Individual
Mean of Dep. Var.	4.69	.14	4.68	.13
Q-Value	0.071	0.071	0.119	0.302

NOTE: tbd

Table E.12: Effect of Salary Delay on Workers' Effort (Alternative Weighting, Hard Images)

	Most Difficult 10 Percent		Most Difficult 25 Percent		Most Difficult 50 Percent	
	Continuous Index	Binary Index	Continuous Index	Binary Index	Continuous Index	Binary Index
ATE	0.0454 (0.0280) [0.0232]	-0.0001 (0.0002) [0.0003]	0.0457 (0.0241)* [0.0223]*	-0.0000 (0.0001) [0.0001]	0.0392 (0.0239) [0.0155]	0.0000 (0.0002) [0.0002]
Observations (First Stage)	199,276	199,276	496,231	496,231	975,265	975,265
Observations (Imputed)	42,053	42,053	114,484	114,484	230,584	230,584
Individuals	297	297	297	297	297	297
SE Cluster	Individual	Individual	Individual	Individual	Individual	Individual
Mean of Dep. Var.	7.06	.001	7.451	.001	7.836	.002
Q-Value	0.411	0.651	0.096	0.602	0.179	0.469

NOTE: tbd

Table E.13: Effect of Salary Delay on Workers' Effort (Alternative Indices and alternative Weighting, Hard Images)

	Most Difficult 10 Percent		Most Difficult 25 Percent		Most Difficult 50 Percent	
	Continuous Index	Binary Index	Continuous Index	Binary Index	Continuous Index	Binary Index
ATE	0.0149 (0.0102) [0.0082]	-0.0002 (0.0003) [0.0003]	0.0186 (0.0070)*** [0.0056]***	-0.0002 (0.0002) [0.0002]	0.0174 (0.0072)** [0.0043]**	-0.0003 (0.0002) [0.0002]
Observations (First Stage)	199,276	199,276	496,231	496,231	975,265	975,265
Observations (Imputed)	42,053	42,053	114,484	114,484	230,584	230,584
Individuals	297	297	297	297	297	297
SE Cluster	Individual	Individual	Individual	Individual	Individual	Individual
Mean of Dep. Var.	3.652	.001	3.831	.001	4.168	.003
Q-Value	0.411	0.411	0.032	0.249	0.064	0.253

NOTE: tbd

E.4 Treatment Saturation

This section shows that results are not meaningfully affected by the level of treatment saturation within a work location.

Table E.14: Treatment Saturation

	Continuous Index		Binary Index	
	Full Treatment Saturation	Partial Treatment Saturation	Full Treatment Saturation	Partial Treatment Saturation
ATE (Pay Cycle One)	0.015 (0.020)	0.019 (0.016)	0.006 (0.004)	0.001 (0.003)
Observations	1,607,926	1,665,507	1,607,926	1,665,507
SE Cluster	Individual	Individual	Individual	Individual
P-Value	0.830		0.210	

NOTE: tbd

E.5 Treatment Effects by Image Difficulty

In this subsection, we examine whether treatment effects are larger in magnitude when the sample is restricted to the most difficult images. To identify difficult images, we consider only untreated observations and apply our usual restriction of discarding the first 500 images for each employee, given the steep initial learning curve. We then calculate the average number of correct labels for each image. Because this sub-analysis was not pre-registered and multiple cutoffs are plausible, we report estimation results for the most difficult 10 percent, 25 percent, and 50 percent of images. We define the most difficult images based on labeling performance among untreated individuals only. For this subset, we compute the average performance for each image and rank images accordingly. The most difficult 10 percent are defined as the lowest-ranked 10 percent of images, and we proceed analogously for alternative cutoffs.


Table E.15 presents the results, showing that treatment effects on the continuous index are larger for the more difficult images. The coefficient of 0.0539 in column 1, estimated on the most difficult 10 percent, indicates that treated employees achieved 0.0539 additional correct labels per image. Relative to the untreated baseline of 7.06 correct labels, this corresponds to an improvement of 0.76 percent. Coefficients in columns 3 and 5 can be interpreted in the same way and correspond to treatment effects of 0.73 and 0.57 percent, respectively. For reference, the ATE using all images was 0.53 percent.

Coefficients on the binary index are all positive but statistically insignificant. However, this reflects the very low share of flawlessly labeled images: only one percent in the most difficult 10 and 25 percent, and just two percent in the most difficult 50 percent. With almost everyone making at least one mistake, it is unsurprising that treatment effects on the binary index cannot be detected statistically.

Table E.15: Treatment Effects on Worker Effort for Difficult Images

	Most Difficult 10 Percent		Most Difficult 25 Percent		Most Difficult 50 Percent	
	Continuous Index	Binary Index	Continuous Index	Binary Index	Continuous Index	Binary Index
ATE	0.0539 (0.0230)** [0.0272]**	0.0002 (0.0003) [0.0004]	0.0547 (0.0205)** [0.0271]**	0.0001 (0.0001) [0.0002]	0.0445 (0.0197)** [0.0305]**	0.0001 (0.0002) [0.0002]
Observations (First Stage)	199,276	199,276	496,231	496,231	975,265	975,265
Observations (Imputed)	42,053	42,053	114,484	114,484	230,584	230,584
Individuals	297	297	297	297	297	297
SE Cluster	Individual	Individual	Individual	Individual	Individual	Individual
Mean of Dep. Var.	7.06	.001	7.451	.001	7.836	.002
Q-Value	0.084	0.479	0.012	0.224	0.037	0.407

NOTE: Tbd

Figure E.5: Treatment Effects on Worker Effort for Difficult Images

NOTE: This Figure shows workers equilibrium effort and firms policies.

E.6 Heterogeneity of Treatment Effects


Table E.16: Average Treatment Effects of Salary Delays by Outside Option

	Performance		Schooling		Gender		WTP	
	Continuous Index	Binary Index	Continuous Index	Binary Index	Continuous Index	Binary Index	Continuous Index	Binary Index
Low Outside Option	0.078*** (0.023)	0.011*** (0.003)	0.090*** (0.027)	0.014*** (0.004)	0.053** (0.021)	0.008** (0.003)		
High Outside Option	-0.005 (0.011)	0.000 (0.003)	0.027 (0.017)	0.002 (0.003)	0.038** (0.018)	0.004 (0.003)		
Low WTP							0.036* (0.021)	0.003 (0.003)
High WTP							0.057*** (0.017)	0.012*** (0.003)
Observations (First Stage)								
Observations (Imputed)								
Individuals								
SE Cluster	Individual	Individual	Individual	Individual	Individual	Individual	Individual	Individual
Mean of Dep. Var.								
p-value (Low=High)	0	.013	.034	.004	.581	.291	.397	.01


NOTE: tbd

Figure E.6: Treatment Effect Heterogeneity (Binary Index)


(a) Heterogeneity by Pre-Treatment Performance



(b) Heterogeneity by Schooling



(c) Heterogeneity by Gender



(d) Heterogeneity by WTP for Salary Certainty



NOTE: This figure illustrates treatment effect heterogeneity. Panel (a) presents heterogeneity by pre-treatment performance, splitting the sample into above- and below-median performers. Panel (b) shows heterogeneity by post-secondary schooling status. Panel (c) reports heterogeneity by gender, and Panel (d) illustrates heterogeneity by WTP for higher salary certainty.

F Theory Appendix

Environment (Formal Description). Time t is discrete and denotes pay cycles. The economy is populated with risk-averse workers and risk-neutral firms, each employing at most one worker. Firms maximize the present value of profits and are permanently characterized by a type $\theta = (A, r)$ drawn from F_Θ , where A denotes productivity and r the interest rate (cost of borrowing). Workers are of type ω and have outside option V^{out} . Firms offer workers a common contractual wage $w > 0$, but each pay cycle decide what fraction $\chi_t \in [0, 1]$ of the owed amount to actually pay. At the start of pay cycle t the firm-worker pair is in state $(B_t, D_t, C_t) \in \mathbb{R}_+^3$, where B_t denote wage arrears owed to a worker, D_t a firm's debt stock, and C_t cash on hand (cash carries zero return). Nature draws a liquidity shock ζ_t , reducing the firm's cash on hand to $C_t^{\text{eff}} = C_t - \zeta_t$. The firm then chooses and pays $\chi_t(B_t + w)$ and services mandatory debt amortization δD_t , using cash C_t^{eff} first and, if needed, new borrowing ℓ_t with $0 \leq \ell_t \leq L_t$. At the end of pay cycle t , any unpaid portion of the contractual obligation becomes next period's arrears. Thus, post-payment arrears carried into cycle $t + 1$ are defined as $B_{t+1} = (1 - \chi_t)(B_t + w)$. If $B_{t+1} > 0$ the firm incurs a reputational (or moral) fixed cost m of withholding wages. The worker observes B_{t+1} and chooses effort η_t , incurring a convex cost $\psi(\eta_t)$, to maximize utility given B_{t+1} and beliefs about next period's payment χ_{t+1} . Continuation beliefs depend only on B_{t+1} . Output $y_t = Af(\eta_t)$ is then realized, and both revenue y_t and any remaining cash C_t^{res} are carried forward as C_{t+1} . The debt stock evolves as D_{t+1} reflecting new borrowing ℓ_t and accrued interest. The employee-employer match survives to the next cycle with probability $p(\eta_t)$. If the match breaks, the firm faces a repayment obligation of outstanding wages, occurring with probability ϕ and scaled by a penalty factor $\xi > 1$. The firm is matched with an identical worker in the next period after paying hiring cost H . All parties apply a common discount factor β .

Workers Beliefs.

Assumption M1 (Workers' Expectation). *Liquidity ζ_t is i.i.d. and independent of firm*

type θ . Observing χ_t does not lead workers to update a posterior over θ ; continuation beliefs depend only on $B_{t+1} = (1 - \chi_t)(B_t + w)$. Aggregation is written as the conditional expectation:

$$\mathbb{E} \left[u(\chi_{t+1}^*(B_{t+1}, D_{t+1}, C_{t+1}; \Theta, \zeta_{t+1}) (B_{t+1} + w) | B_{t+1}) \right]. \quad (14)$$

where the expectation is taken over (Θ, D, C, ζ) under a product measure and impose distributional assumptions for tractability (details in Appendix XXX).

The Firm Optimization. Per-period firm profits are defined as output minus all potential costs the firm has to pay (debt amortization, wages and the cost of wage withholding). Formally, this is defined as

$$\pi_t(B_t, D_t, C_t, \chi_t; \theta) = y_t - \delta D_t - \chi_t(B_t + w) - m(B_{t+1}). \quad (15)$$

Given the worker's response $\eta_t^*(B_{t+1})$, a firm of type $\theta = (A, r)$ solves

$$\begin{aligned} V_f(B_t, D_t, C_t; \theta) = \mathbb{E}_\zeta \left[\max_{\chi_t \in [0,1]} \left\{ \pi_t(B_t, D_t, C_t, \chi_t; \theta) \right. \right. \\ \quad \left. \left. + \overbrace{\beta p(\eta_t^*(B_{t+1})) V_f(B_{t+1}, D_{t+1}, C_{t+1}; \theta)}^{\text{Continuation value if match survives}} \right. \right. \\ \quad \left. \left. + \underbrace{(1 - p(\eta_t^*(B_{t+1})))}_{\text{Probability match breaks}} \underbrace{(\phi \xi B_{t+1} - H)}_{\text{Expected fine}} \underbrace{-}_{\text{Rehiring cost}} + \overbrace{\beta V_f(0, D_{t+1}, C_{t+1}; \theta)}^{\text{Value after rehiring}} \right\} \right] \end{aligned} \quad (16a)$$

subject to

$$\pi_t(B_t, D_t, C_t, \chi_t; \theta) = y_t - \delta D_t - \chi_t(B_t + w) - m(B_{t+1}) \quad (16b)$$

$$C_t^{\text{eff}} + \ell_t = \delta D_t + \chi_t(B_t + w) \quad (16c)$$

$$0 \leq \ell_t \leq L \quad (16d)$$

$$D_{t+1} = (1 + r)(1 - \delta)D_t + \ell_t \quad (16e)$$

$$0 \leq D_{t+1} \leq D^{\text{max}} \quad (16f)$$

$$C_t^{\text{res}} = \max\{0, C_t^{\text{eff}} - \delta D_t - \chi_t(B_t + w)\} \quad (16g)$$

$$C_{t+1} = y_t + C_t^{\text{res}} \quad (16h)$$

The firm's value function is a discontinuous function because ℓ switches on at $\ell = 0$, and the borrowing cap may bind at $\ell = L$. Additionally, $m(B)$ also jumps at $B = 0$. The function is piecewise smooth in χ and optimal χ^* may be at corners or at piecewise interior points.

Equilibrium. Parameters (L, δ, r) and the support of ζ ensure: (i) for each (B, D, C, θ) there exists a feasible $\chi \in [0, 1]$ via (??); (ii) under optimal policies, D and C remain bounded by (??); (iii) u is continuous and strictly concave, p is continuous, increasing and strictly concave, $m(\cdot)$ is bounded and weakly increasing, and F_Θ, F_ζ have bounded support.


Definition 1 (Stationary equilibrium). *A stationary equilibrium consists of $(V_f, V_w, \chi^*, \eta^*, U_{\text{pay}})$ such that:*

- (i) **Worker optimality:** *Given U_{pay}, V_w solves (??), and η^* satisfies (3).*
- (ii) **Firm optimality:** *Given η^* , for each θ , $\chi^*(\cdot)$ solves (16) subject to (16b)-(16h).*
- (iii) **Aggregation consistency:** *U_{pay} is generated by χ^* via (14).*

We solve for the equilibrium numerically using value function iteration, with computational details provided in Appendix XXX. Figure F.7 illustrates the resulting equilibrium, showing

optimal worker effort and aggregate firm policies for a parameterization calibrated to approximate the economic environment of Nigeria. Given reasonable low enforcement, we see that optimal effort can increase initially for low levels of arrears. At the same time, firms find it optimal not to pay workers with the repayment share declining as arrears grow.

Figure F.7: Worker Effort and Firm Policy in Equilibrium



NOTE: This Figure shows workers equilibrium effort and firms policies.

G Supplementary Data Collections and Interventions

This section describes the supplementary data collections and interventions.

Weekly Employee Surveys (Employee Sample). To illustrate the effects of wage withholding on employees' self-reported wellbeing, we conducted weekly employee surveys via WhatsApp. Every Friday, employees received a brief WhatsApp questionnaire, framed as originating from a third-party recruitment agency interested in learning about their experiences with this particular employer. This framing encouraged candid responses regarding salary payments and working conditions. Employees were incentivized to complete the survey through a small monetary reward and the opportunity to win a larger cash prize, resulting in a relatively high response rate of 34 percent. Summary statistics and further methodological details are provided in Appendix XXX. We use these data in Section XXX, complementing our primary experimental results with information on employees' subjective experiences.

G.1 End-of-Job Survey (Employee Sample)


To elicit employees' beliefs about salary delays and their general perceptions of the job, we conducted an end-of-employment survey on the final workday of both employment rounds. The first survey was administered on May 2, 2025, the last workday of the first experimental employment round, and the second on October 3, 2025, the last workday of the second round. Both surveys were administered prior to the debriefing. We surveyed all employees present on those days—219 in the first round and 268 in the second. Surveys were administered individually to prevent responses from being influenced by peers, for example through social image concerns.

First, we asked employees about the purpose of the job. This was included to validate the experimental design and to ensure that employees perceived the job as genuine and primarily about image labeling. Responses are reported in Appendix Figure G.8 for both employment rounds. The figure shows that the vast majority of employees (80.4 percent in the first round

and 78.4 percent in the second) believed the primary purpose of the job was image labeling. The remaining 19.6 percent and 21.6 percent, respectively, believed the job served another purpose, but only 4.6 percent and 2.2 percent, respectively, explicitly stated that it was to assess the productivity effects of salary delays. We interpret these results as evidence that our experimental setup was realistic and perceived as genuine.


Second, we asked employees who experienced salary delays to self-assess how they responded to those delays. This was a multiple-choice question, and employees could select one or more options (conflicting responses were not possible). Results are reported in Appendix Figure G.9. Bars 1 and 2 refer to self-assessments of work effort: 31 percent of employees reported working harder when experiencing salary delays, while only 2 percent reported working less hard. Bars 3 to 5 refer to attendance: most employees (73 percent) indicated that they attended work as usual when their salaries were delayed, while small fractions reported attending more or less, respectively.

Figure G.8: Perceived Purpose of Employment



NOTE: This figure presents results from the end-of-employment survey conducted for both experimental employment rounds. The survey was administered to all 219 employees attending the final workday of the first employment round and all 268 employees attending the final workday of the second round. The figure shows employees' single-choice responses to the question: "In your opinion, what was the main purpose of this job and the program?"

Figure G.9: Self-Assessment of Responses to Salary Delays



NOTE: This figure presents results from an end-of-employment survey administered to 219 employees on the final workday. It shows employees' multiple-choice responses (conflicting responses were not possible) to the question: "How did you respond when your salary was delayed?"

G.2 Employment Terms Perception Survey (Recruitee Sample)

We conducted a follow-up survey with recruitees who were not hired to elicit their perceptions of the employment terms across the different treatment conditions. The survey assessed how treatments influenced recruitees' concerns about salary uncertainty. To avoid priming respondents about this issue, we framed the survey as a routine follow-up by the recruitment agency to evaluate the performance of the interviewer. A brief questionnaire was administered either by phone (by a different enumerator) or through automated WhatsApp messages. It included questions about recruitees' initial interview or information session, followed by their evaluation of the employment terms. We surveyed 204 individuals in total: 125 via WhatsApp and 79 via phone. Participants received a small monetary incentive in the form of mobile phone recharge cards for completing the survey. The administered questionnaire was identical between the phone and WhatsApp surveys. Respondents were asked how concerned they would be about receiving their salary *on time* and *at all* if they were to start

the job. They could choose among four options: “Very concerned,” “Concerned,” “Little concerned,” and “Not concerned,” corresponding to a four-point Likert scale. Participants were asked how concerned they would be about receiving the salary on time or at all if they were to start the job. Participants had four response options *Very Concerned*, *Concerned*, *Little Concerned*, and *Not Concerned* corresponding to a 4-point Likert scale.


We then estimate the following regression:

$$y_i = \alpha + \beta_2 Z_{i2} + \beta_3 Z_{i3} + \varepsilon_i. \quad (17)$$

where y_i is the numeric value of individual i ’s response, and Z_{ij} is an indicator for the type of employment terms j corresponding to one of the three treatment arms given to individual i . Treatment Arm 1 (Control Arm) serves as the omitted reference category.

Appendix Figure G.10 presents the survey results. The figure illustrates average levels of recruits’ concern about receiving their salary on time and, respectively, about receiving it at all. The patterns are fairly similar: individuals in the control arm report concern levels slightly above “Little Concerned.” Those assigned to the uncertainty arm exhibit nearly identical levels of concern, which are not statistically distinguishable from the control group. This provides additional evidence of the high prevalence of wage withholding: telling people that salary may not be paid does not significantly change their concern about receiving their salary on time or at all. In contrast, the certainty treatment reduces concern. Recruits offered employment terms with guaranteed salary are significantly less concerned than those in the control arm, reporting average levels of concern slightly below “Little Concerned.”

Figure G.10: Salary Concerns by Treatment Arm (Follow-Up Survey)




NOTE: This figure presents results from the follow-up survey with recruits on their perceptions of the treatment arms, estimated using Equation (17). Reported values are averages on a four-point Likert scale (1 = Very concerned, 2 = Concerned, 3 = Little concerned, 4 = Not concerned). We display 95% confidence intervals around the averages, based on heteroskedasticity-robust standard errors.

H Selection Robustness

I Complier Characteristics

Figure I.11: Complier Characteristics



NOTE: This figure shows characteristics of always takers, compliers, and never takers.

Table I.17: First-Stage for Compliers

	(1)	(2)
Salary Guarantee	0.13*** (0.03)	0.08** (0.03)
Observations	1,079	1,079
R-Square	0.45	0.01
F-Statistic	15.81	6.12
Date Control	Yes	No
Salary Control	Yes	No

NOTE: tbd

J Choice Experiment

J.1 Design and Implementation

The choice experiment elicited recruits’ preferences over different job offers and was administered during the job interview for the sample responding to the job advertisement, respectively during the job information session for the in-person recruited sample. Recruits were shown one of three choice blocks, each containing four binary choices between job offers. The job offers varied in two attributes: salary, and information on payment modalities. Salary was randomly varied across offers within the range of 55,000 to 85,000 NGN, while payment information varied according to the three treatment arms described in Section 4.1. Additionally, all job descriptions included the same information on location, hours, and job type.

Jobseekers were asked to indicate which of the shown job offers they would accept; possible responses included accepting one, both, or neither. We used a Bayesian D-efficient algorithm — a modified Fedorov algorithm (Cook and Nachtsheim, 1980; Zwerina et al., 1996; Carlsson and Martinsson, 2003) — implemented with the software developed by Hole (2015), to select

the first three choices in each of the three choice blocks to maximize the statistical power of the design. Only the first three choices were used for estimating preferences, as specified in the pre-analysis plan. The fourth choice in each block was relevant for the branch of the experiment in which one of the respondent’s choices could be implemented as an actual job offer, as described in the next paragraph.

Appendix Figure J.1 shows two example choices presented to recruits. Panel (a) illustrates a choice between a job offer corresponding to treatment arm 3 (salary certainty arm) and treatment arm 2 (uncertainty arm). Panel (b) shows a choice between treatment arm 1 (control arm) and treatment arm 2 (uncertainty arm).

Most employees were told that their choices would help us (the recruitment agency) to offer job opportunities matching their preferences. This constituted an incentive-compatible design for eliciting truthful responses. To further strengthen incentive compatibility, a subset of recruits was randomly assigned to a different experimental condition in which they were told — truthfully — that one of their selected choices would be implemented as an actual job offer. While this approach is highly effective for eliciting truthful responses, the D-efficient design of the choice experiment limited the range of job-offer combinations we could feasibly implement. For this reason, we included a fourth choice in each block, which was not used in the preference estimation but provided additional flexibility for making offers in this subset of cases. The set of job offers we could extend remained constrained, which is why we implemented this highly incentive-compatible condition only for a subset of recruitees.

J.2 Theory and Estimation

To understand how employees value different job attributes and estimate WTP for higher salary certainty, we specify a discrete choice model.⁵⁹ Following the conceptual setup of McFadden (1974), we model individuals’ utility from job choice j , denoted U_j , as deterministically dependent on observable job characteristics x_j and the w_j . Utility additionally

⁵⁹In this case, WTP is defined as the willingness to give up salary for greater salary certainty.

depends stochastically on the unobservable term ε_j , so that utility U_j from job choice j can formally be expressed as

$$U_j = v(x_j, w_j; \beta) + \varepsilon_j, \quad (18)$$

where β is a vector of parameters. While utility also depends on the stochastic component, ε_j , it is deterministic from the perspective of the individual making the job choice. We further assume that individuals choose their utility maximizing job option j from the choice set C . The probability that an individual then chooses job j from the choice set C can be written as

$$\Pr(j|C) = \Pr(U_j > U_i) = \Pr(v_j + \varepsilon_j > v_i + \varepsilon_i) = \Pr(v_j - v_i > \varepsilon_i - \varepsilon_j) \quad \forall i \in C. \quad (19)$$

To estimate this choice probability, we impose the standard restrictions on the structure of the utility function. First, we assume that ε_j follows a Type I extreme value distribution. Second, we assume that the utility function is linear in parameters. Third, we rely on the independence of irrelevant alternatives (IIA) assumption, which requires that the relative probability of choosing job j over i is unaffected by the inclusion or exclusion of other alternatives in the choice set. Under these assumptions the probability of choosing job j from the choice set C can be estimated using a conditional logit model. The probability of choosing job j as a function of the job characteristics x_l with $l = 1, \dots, m$ and the salary w can then be written as

$$\Pr(j) = \frac{\exp(\sum_{l=1}^m \beta_l x_{jl} + \beta_w w_j)}{\sum_{j \in C} \exp(\sum_{l=1}^m \beta_l x_{jl} + \beta_w w_j)}. \quad (20)$$

We estimate the choice probabilities in equation (20) via maximum likelihood. Salary w is a continuous variable ranging from 55,000NGN to 85,000NGN. We include two additional variables capturing job attributes l . First, salary delay d is a continuous variable of the specified probability of salary delay between 0% and 50% (0% is the control condition). Second, the binary variable $s \in \{0, 1\}$ is a categorical variable indicating the usage of an automated payment system by the firm which is also reflected in the job offer. Standard errors are clustered at the individual level.

An advantage of the conditional logit model is that marginal rates of substitution (MRS) are straightforward to compute. This is particularly relevant in this case, because it allows us to derive a valuation of the different job attributes relative to salary. We interpret the MRS between a job characteristic and salary as a WTP for that job characteristic. The MRS between two attributes in the conditional logit model is simply the ratio of their estimated coefficients. Accordingly, the WTP for job characteristic l is given by:

$$WTP_l = \frac{\partial U / \partial x_l}{\partial U / \partial w} = \frac{\beta_l}{\beta_w} \quad (21)$$

J.3 Results

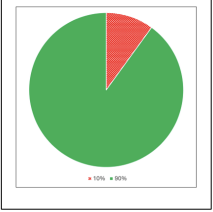
Table J.2 shows the full sample estimation results for the WTP estimates. Table J.3 replicates Table J.2 using a subsample only whose choice was implemented, making their answers highly incentive compatible. The tables show the MRS, i.e. coefficient ratio, between a job offer that offers higher salary certainty through the automated payment system and salary. The

table also shows the MRS between the specified probability that salary delay may occur for a given employee and salary. Column (1) shows the estimation results using the entire sample. Column (2) shows the estimation results using the sample that responded to the job advertisement only, column (3) uses the in-person recruited and non-incentivized sample only while column (4) shows the estimation results for in-person recruited sample that had to be incentivized to participate in the job information session. The coefficient on ‘WTP for Salary Certainty’ in column (1) implies that jobseekers were willing to give up 27,223 NGN (~ 18 USD) in monthly salary for a job offer that included the automated payment system and conveyed higher salary certainty. This is a substantial amount corresponding to 39% of the minimum wage of 70,000 NGN.

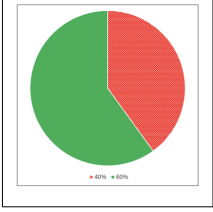
The coefficient on ‘WTP for Delay’ in column (1) implies that jobseekers are willing to accept a 1% higher probability that their salary will be delayed for a salary increase of 733 NGN.

Figure J.1: Examples of Job Offer Choices in the Choice Experiment

(a) Example Choice One

	Job Offer A	Job Offer B
Job Type	Image classification	Image classification
Location, hours	Satellite Town, 9:00am – 5:00pm	Satellite Town, 9:00am – 5:00pm
Monthly Salary	55,000 NGN	85,000 NGN
Payment	Biweekly This company offers an automated payment system to guarantee on-time payment: your salary would be transferred automatically from a bank account with enough money to cover the salary payments. No worker who received this has reported any issues with their salary.	Biweekly Things in Nigeria are difficult at the moment, also for this company: every month there is a 10% chance your salary will not be paid 

(b) Example Choice Two

	Job Offer E	Job Offer F
Job Type	Image classification	Image classification
Location, hours	Satellite Town, 9:00am – 5:00pm	Satellite Town, 9:00am – 5:00pm
Monthly Salary	55,000 NGN	85,000 NGN
Payment	Biweekly	Biweekly Things in Nigeria are difficult at the moment, also for this company: every month there is a 40% chance your salary will not be paid 

NOTE: This figure shows two example choices from the choice experiment as presented to recruits.

Table J.1: Attributes and Variation of Job Offers in the Choice Experiment

Job Attribute	Attribute Levels
Salary	55,000 NGN 60,000 NGN 65,000 NGN 70,000 NGN 75,000 NGN 80,000 NGN 85,000 NGN
Payment Mopdalitie	Biweekly Payments Biweekly Payments This company offers an automated payment system to guarantee on-time payment: your salary would be transferred automatically from a bank account with enough money to cover the salary payments. No worker who received this has reported any issues with their salary. Biweekly Payments Things in Nigeria are difficult at the moment, also for this company: every month there is a {10, 20, 30, 40, 50} percent chance your salary will not be paid.
Job Type	Image Classification
Location and Hours	Satellite Town, 9:00am to 5:00pm

NOTE: This table shows the variation in job offers presented in the choice experiment. We vary salary and information on salary certainty according to the three treatment arms described in Section 4.1. All job offers provide the same information on job type, location, and hours.

Table J.2: Choice Experiment: Full Sample

	Entire Sample	Job Ad Sample	In-Person Sample (No Incentive)	In-Person Sample (Incentive)
WTP for Sal. Certainty	26.64*** (0.99)	25.99*** (1.17)	26.31*** (2.33)	28.37*** (2.41)
WTP for Delay	-0.71*** (0.02)	-0.63*** (0.03)	-0.72*** (0.05)	-0.95*** (0.06)
Observations	6,562	3,826	826	1,910
Individuals	1,110	641	139	330
SE Cluster	Individual	Individual	Individual	Individual


NOTE: This table shows the estimates of the choice experiment administered during the job interview. The table shows the coefficient ratios from a conditional Logit estimation, which can be interpreted as marginal rates of substitution. Standard errors are initially clustered at the individual level, and we calculate the standard error of the coefficient ratio using the delta method. Column (1) shows the estimation results using the entire sample. Column (2) shows the estimation results using the sample that responded to the job advertisement only and column (3) uses the in-person recruited sample only. Column (4) shows the estimation results for the sub-sample which was informed that one of their choices in the choice experiment would be implemented as their job offer.

Table J.3: Choice Experiment: Implemented Choices Subsample

	Entire Sample	Job Ad Sample	In-Person Sample (No Incentive)	In-Person Sample (Incentive)
WTP for Sal. Certainty	25.98*** (1.21)	25.11*** (1.43)	25.32*** (1.87)	28.54*** (3.41)
WTP for Delay	-0.70*** (0.03)	-0.62*** (0.04)	-0.72*** (0.05)	-0.95*** (0.08)
Observations	3,468	2,052	440	976
Individuals	586	344	74	168
SE Cluster	Individual	Individual	Individual	Individual

NOTE: This table shows the estimates of the choice experiment administered during the job interview. The table shows the coefficient ratios from a conditional Logit estimation, which can be interpreted as marginal rates of substitution. Standard errors are initially clustered at the individual level, and we calculate the standard error of the coefficient ratio using the delta method. Column (1) shows the estimation results using the entire sample. Column (2) shows the estimation results using the sample that responded to the job advertisement only and column (3) uses the in-person recruited sample only. Column (4) shows the estimation results for the sub-sample which was informed that one of their choices in the choice experiment would be implemented as their job offer.

Figure J.2: Distribution of Individual WTP for Salary Certainty



NOTE: This figure shows the distribution of individual-level WTP values, expressed in USD. These values represent the amount individuals are willing to forgo from their monthly salary in exchange for higher salary certainty.

Table J.4: Workforce Composition Effects

	Continuous Index		Binary Index	
	(1)	(2)	(3)	(4)
In-person recruits	0.025 (0.035)	0.034 (0.032)	0.001 (0.006)	0.001 (0.006)
R-Square	0.0002	0.6288	0.0000	0.7259
Observations	505,274	505,274	505,274	505,274
Individuals	221	221	221	221
SE Cluster	Individual	Individual	Individual	Individual
Date FE	No	Yes	No	Yes
Image FE	No	Yes	No	Yes

NOTE: This table reports the estimation results from equation (11), comparing productivity across three samples: job-ad recruits (sample III and the omitted category), in-person recruits without incentive (sample I), and in-person recruits with incentive (sample II). All employees were hired under contracts with salary guarantees. Standard errors, clustered at the individual level, are reported in parentheses.

Appendix References

- Carlsson, Fredrik and Peter Martinsson**, “Design techniques for stated preference methods in health economics,” *Health Economics*, 2003, 12 (4), 281–294.
- Chen, Angela**, “Millions of Workers Are Training AI Models for Pennies,” *Wired*, 2023. Accessed: September 5, 2025.
- Cook, R. Dennis and Christopher J. Nachtsheim**, “A Comparison of Algorithms for Constructing Exact D-Optimal Designs,” *Technometrics*, 1980, 22 (3), 315–324.
- G2 Learning Hub**, “20+ Data Labeling Statistics You Need to Know,” 2023. Accessed: September 5, 2025.
- Hole, Arne Risa**, “DCREATE: Stata module to create efficient designs for discrete choice experiments,” Statistical Software Components, Boston College Department of Economics Aug 2015.
Humans Are Essential for Training AI
- Humans Are Essential for Training AI***, *Le Monde*, 2024. Accessed: September 5, 2025.
- Kässi, Otto, Vili Lehdonvirta, and Fabian Stephany**, “How Many Online Workers are there in the World? A Data-Driven Assessment,” 2021. Accessed: September 5, 2025.
- Market.us Insights**, “Data Labeling Solution and Services Market News,” 2024. Accessed: September 5, 2025.
- McFadden, Daniel**, “Conditional Logit Analysis of Qualitative Choice Behavior,” 1974.
The Hidden Workforce Behind Artificial Intelligence
- The Hidden Workforce Behind Artificial Intelligence***, *Financial Times*, 2023. Accessed: September 5, 2025.
- Vazquez-Bare, Gonzalo**, “Identification and estimation of spillover effects in randomized experiments,” *Journal of Econometrics*, 2023, 237 (1), 105237.
- Zwerina, Katerina, Joel Huber, and Warren Kuhfeld**, “A General Method for Constructing Efficient Choice Designs,” Working Paper, Fuqua School of Business, Duke University 1996.